



**HDFS**





# HADOOP DISTRIBUTED FILE SYSTEM

# Agenda



- ✓ HFDS Design
- ✓ HDFS Architecture
- ✓ HDFS Storage Daemons
- ✓ Name Node, Data Node
- ✓ Secondary Name Node & Check-pointing
- ✓ HDFS Storage Architecture
- ✓ HDFS File Write & Read
- ✓ Hadoop High Availability
- ✓ Basic HDFS Operations



# HDFS Design



HDFS is a filesystem designed for storing **very large files** with **streaming data access** patterns, running on clusters of **commodity hardware**.

## Very Large Files

100s of GBs or even TBs.

## Streaming Data Access

Built for “Write-once Ready-many-times” pattern. Time to read the whole dataset is more important than the latency in reading the first record.

## Commodity Hardware

Where chance of node failure is high.

# Hadoop Master/Slave Architecture



Name Node



Master Node



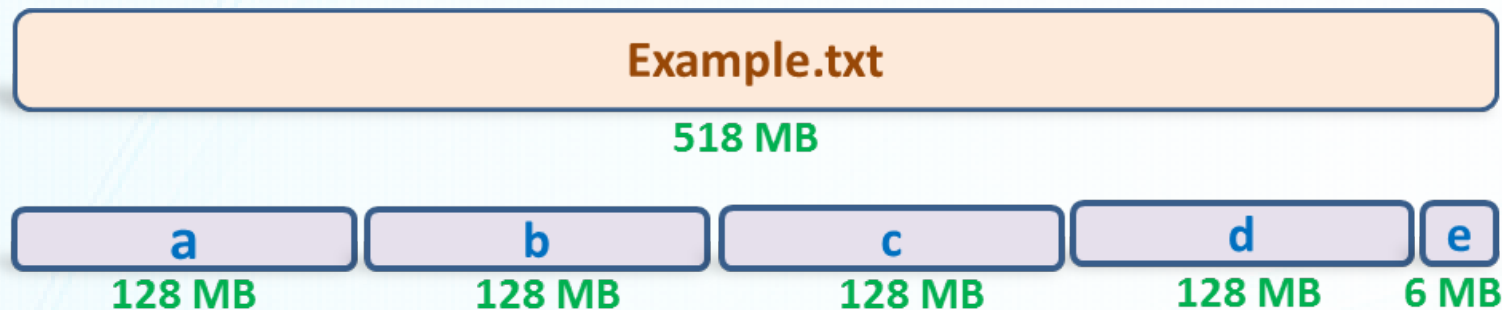
Slave Nodes

Data Nodes (Commodity Hardware)

# Blocks



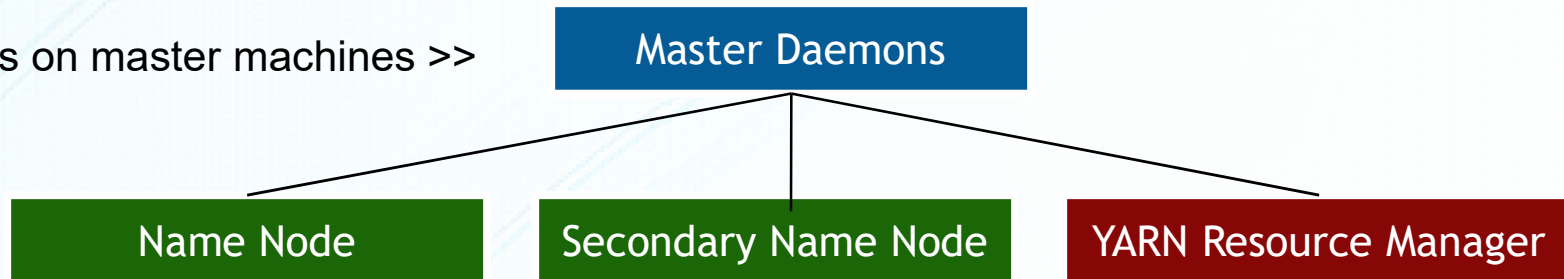
- In HDFS, files are broken into blocks of 128MB and are distributed across various nodes in the cluster.
- Blocks are the basic unit of abstraction in HDFS
- **Advantages:**
  - a file can be larger than any single disk in the network
  - simplifies the storage subsystem
  - fits well with replication for providing fault tolerance and availability



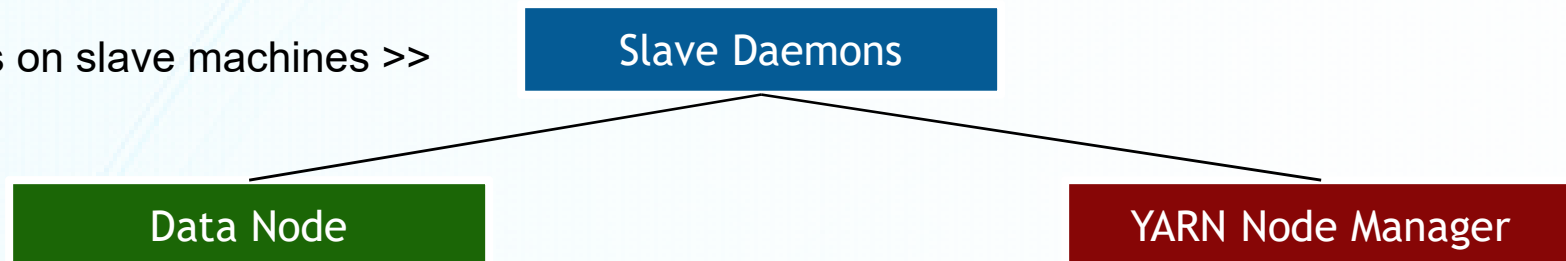
# Hadoop Daemon Processes



Runs on master machines >>



Runs on slave machines >>

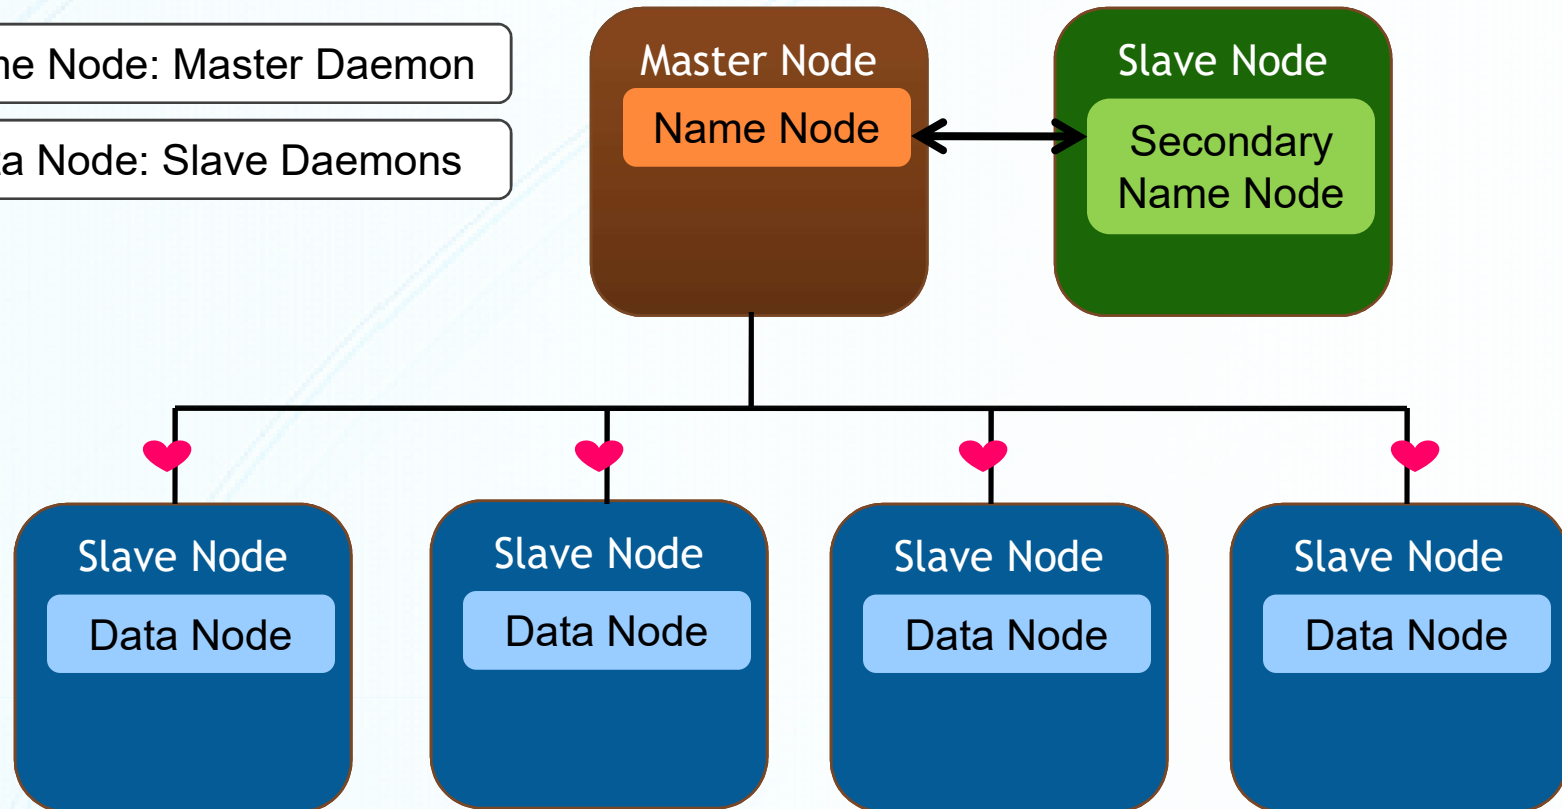


# HDFS Storage Daemons



Name Node: Master Daemon

Data Node: Slave Daemons



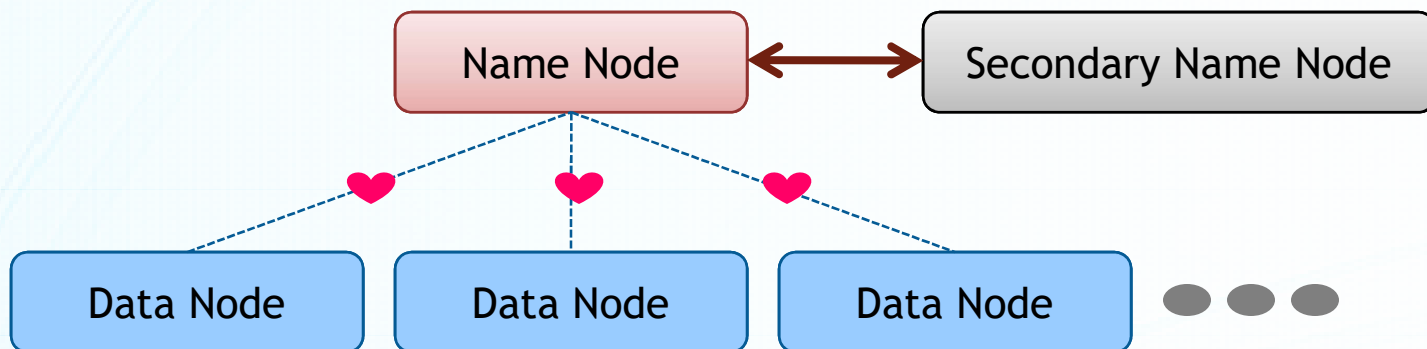
Data Nodes communicate with Name Node through heartbeats ♥



# Name Node



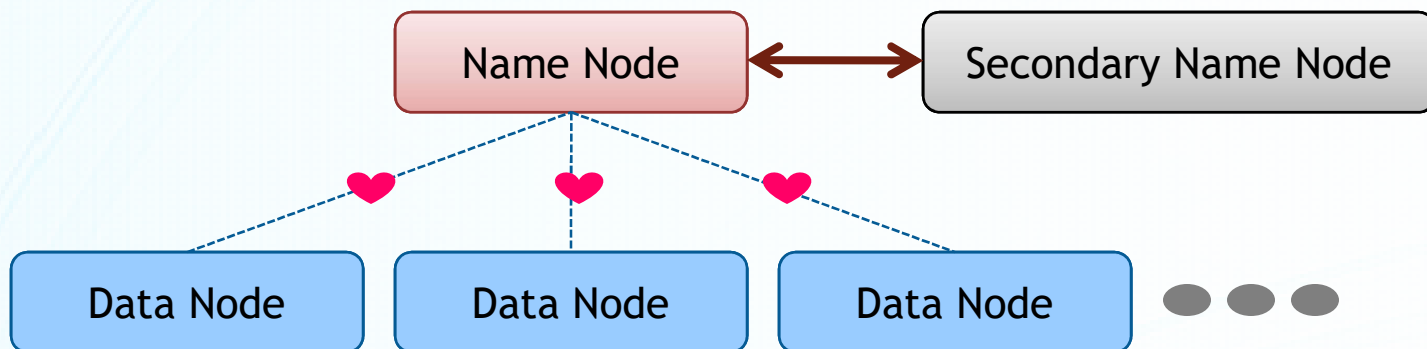
- Name Node is the master daemon that maintains and manages Data Nodes
- Maintains meta data about data blocks such as location of blocks stored, size of the files, permissions, hierarchy etc in a file called **FsImage**.
- Receives heart-beat and block report from all Data Nodes.



# Data Node



- These are slave daemons that runs on each slave machines
- Stores and manages actual data
- Serves Read and Write requests from the client
- Sends heartbeat to Name Node once every 3 sec.



# FsImage & Edit Log Files



## FsImage File

- FsImage file contains the complete state of the file system namespace since the start of the NameNode.
- Each node is an internal representation of a file or directory's metadata and contains such information as the file's replication level, modification and access times, access permissions, block size, and the blocks a file is made up of.
- For directories, the modification time, permissions, and quota metadata is stored.

## EditLog File

- EditLog file is stored in-memory and records all incremental changes happening on HDFS and is periodically merged with FsImage during check-pointing process.

# Secondary Name Node



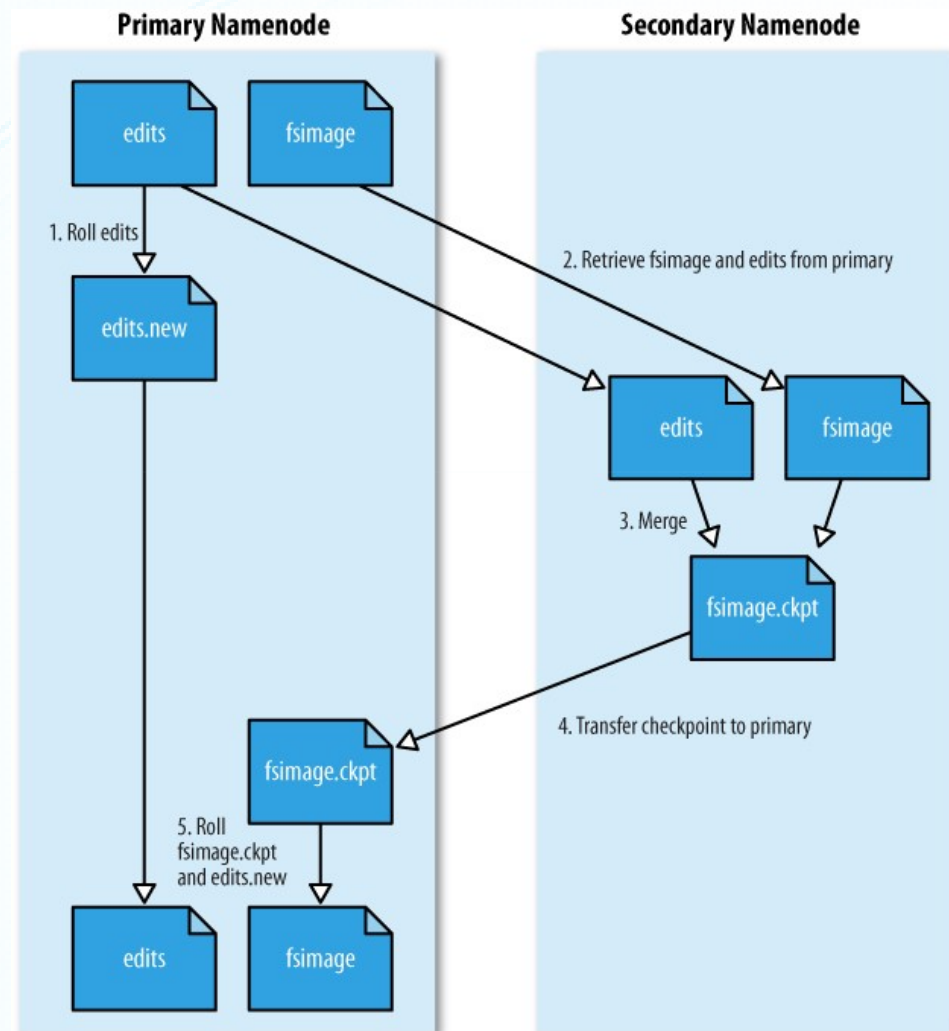
**Check-pointing** is the process of combining edit-logs with FsImage

SNN takes the responsibility of check-pointing, thereby reducing load on Name Node

Allows faster fail-over as it prevents the edit-log from getting too huge.

Check-pointing happens periodically.

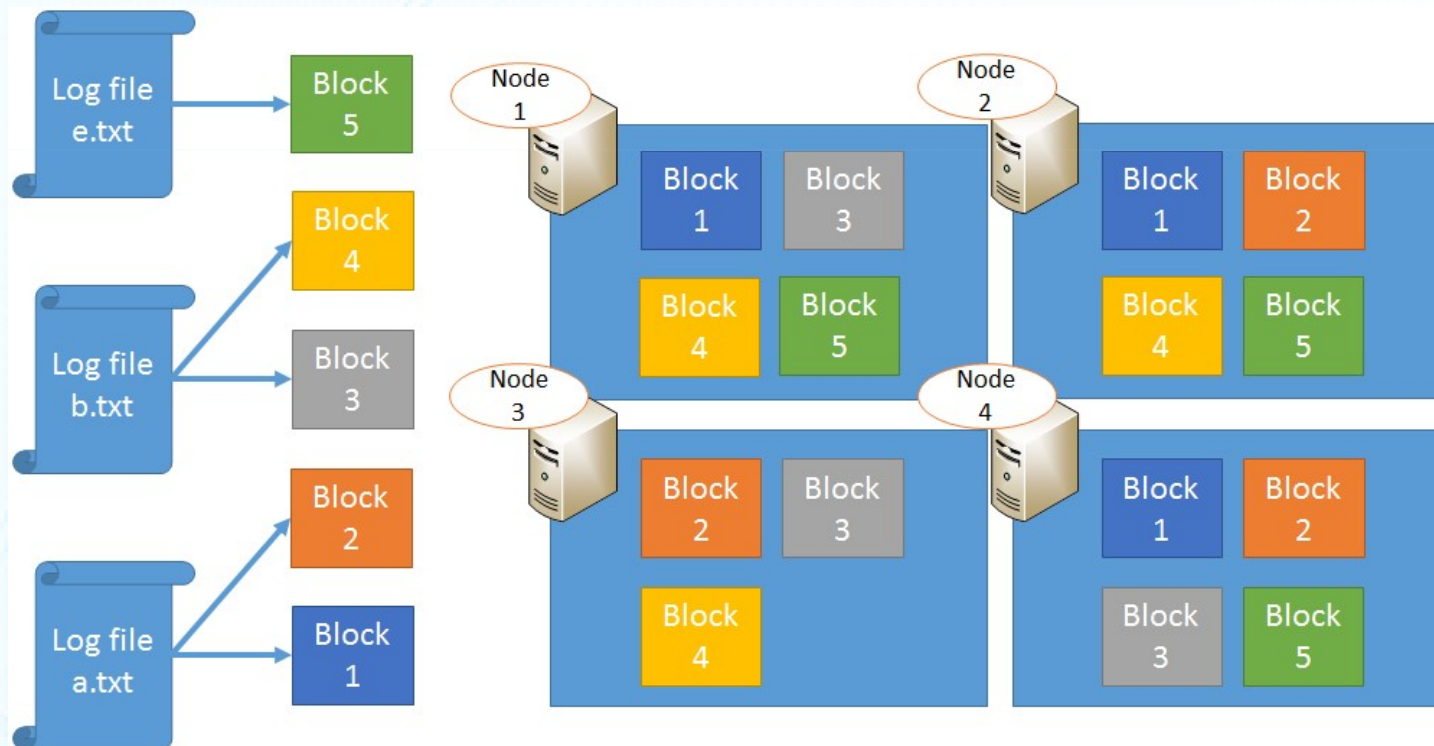
SNN is not a hot stand-by for NN. It just does metadata backup and is used to rebuild a failed NN



**How is data stored in Data Nodes?**

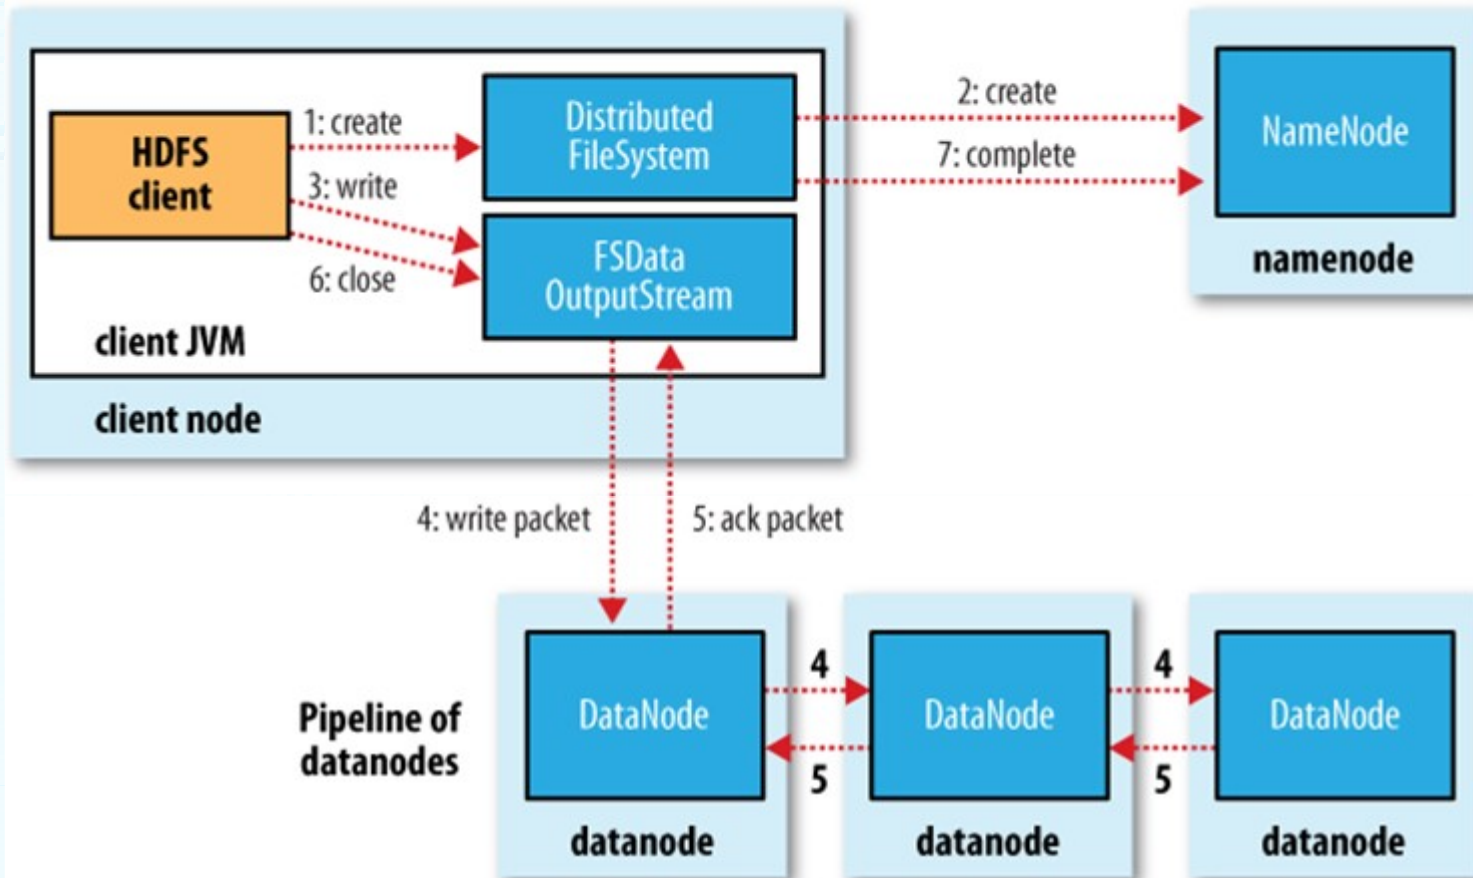
# HFDS Blocks & Replication

- HDFS is a block structured file system.
- Each file is stored on HDFS as blocks of 128MB (default)
- Each block is replicated 3x (default) to deal with Data Node failure



**How are files written to and read from HDFS?**

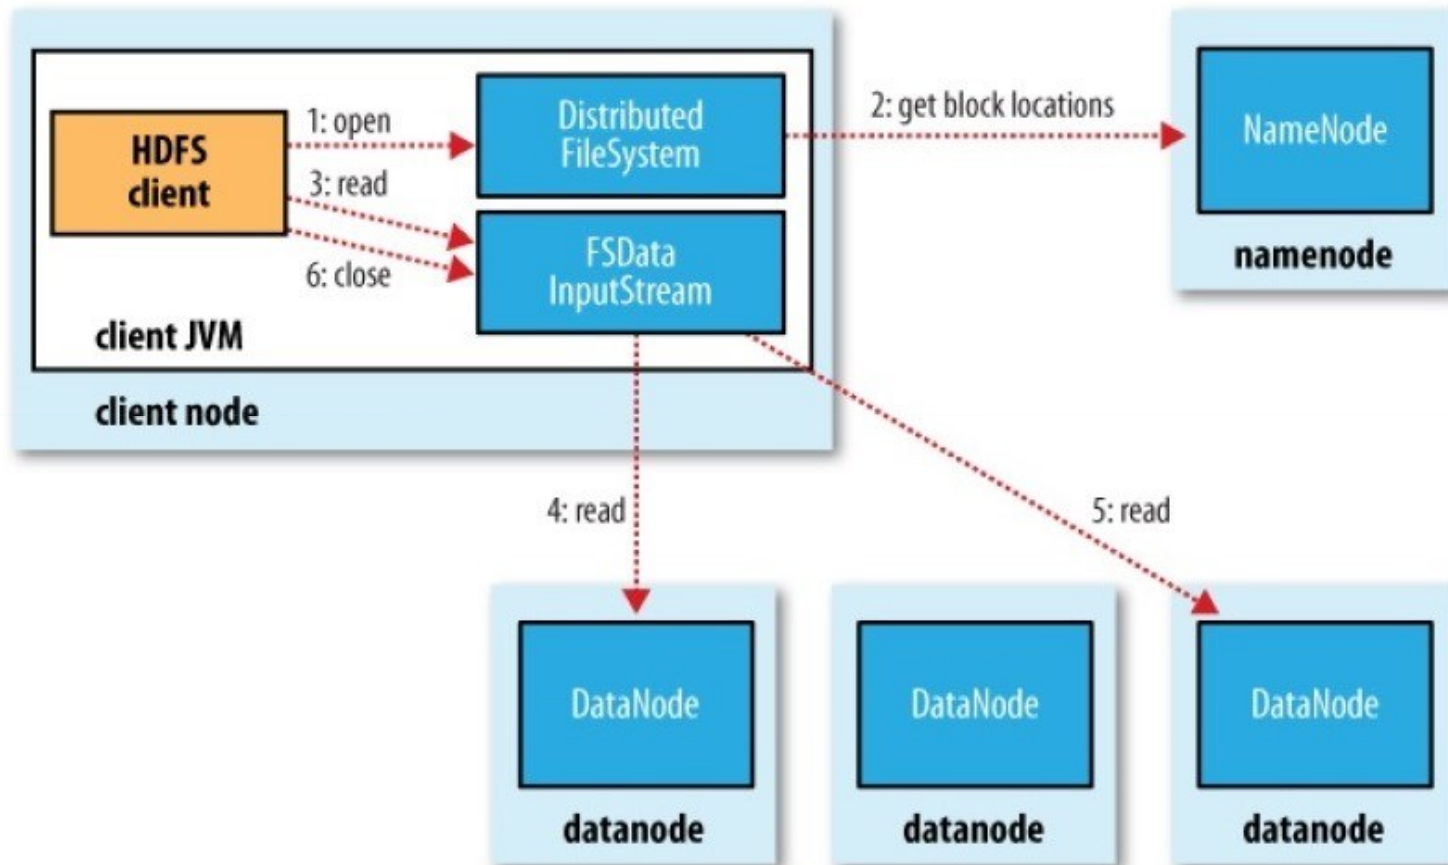
# Anatomy of HDFS File Write



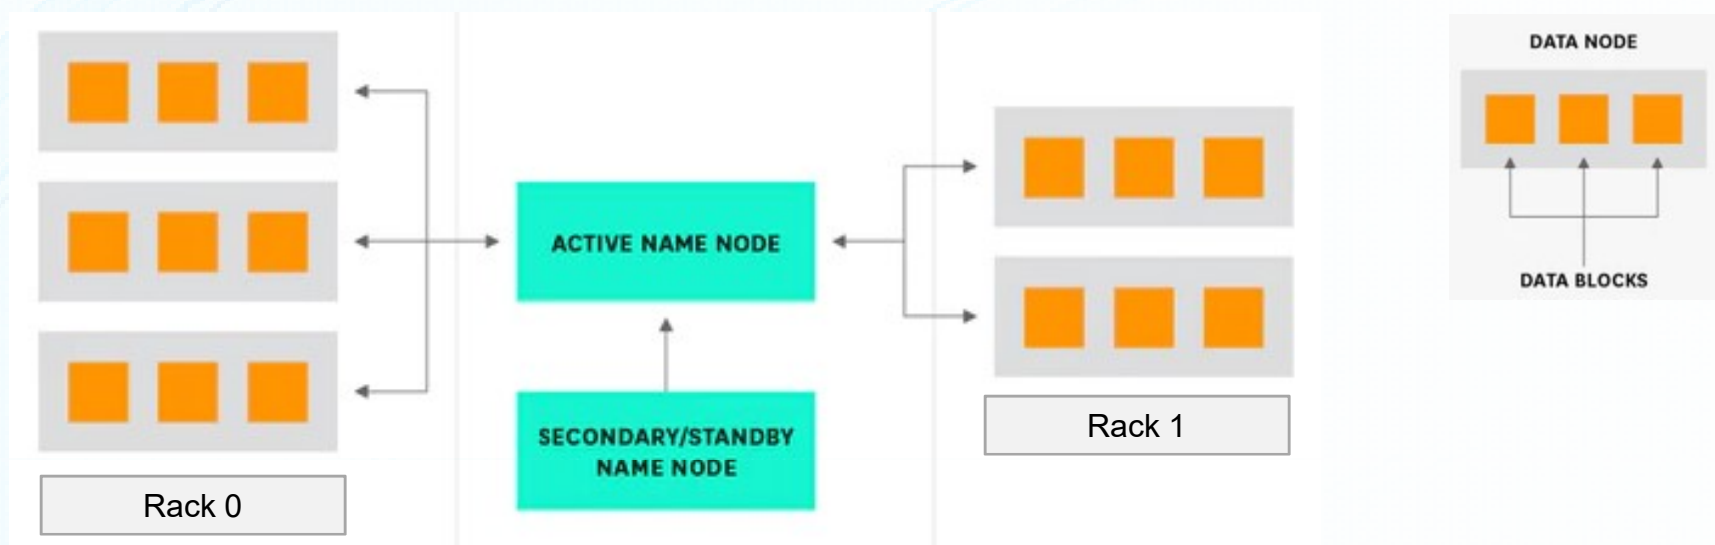
Replications are written sequentially. Blocks are written in parallel.



# Anatomy of HDFS File Read

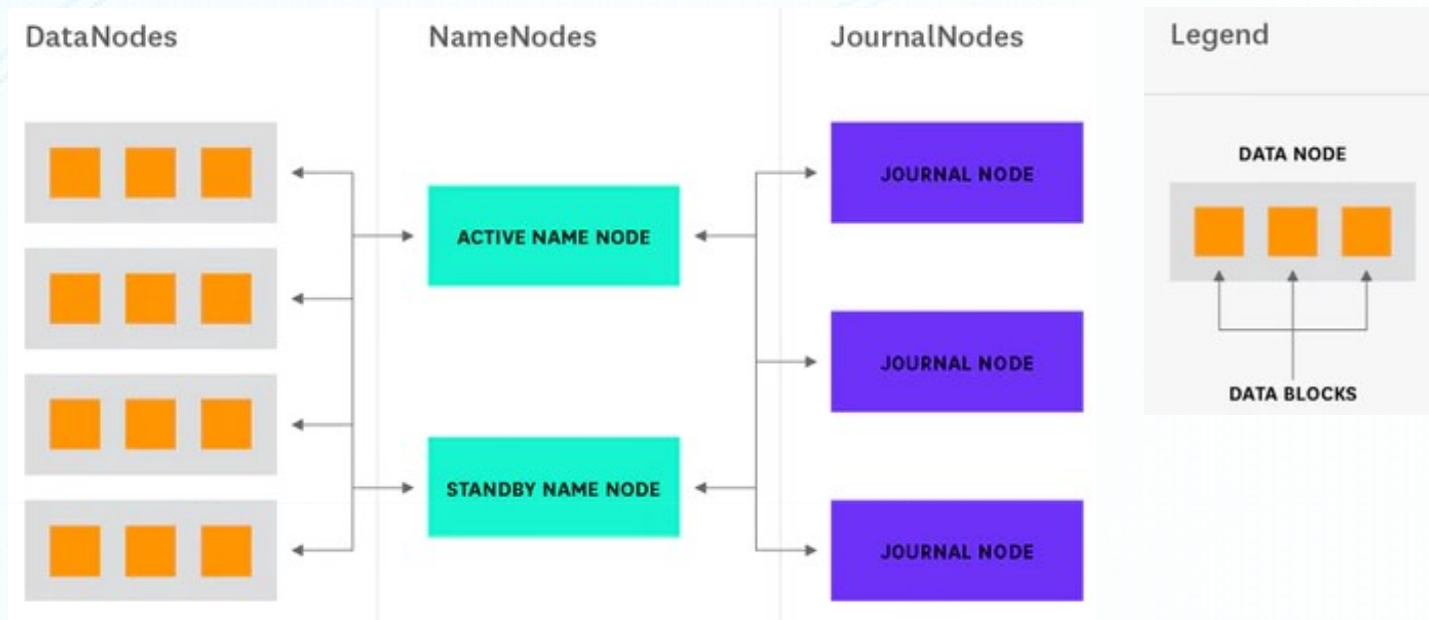


# Hadoop Cluster Architecture with out HA



- In the early versions of Hadoop, NameNode was the 'single point of failure'.
- SNN does not provide an alternative to NN, but only helps to quickly recover from NN failures. However, if NN is down, the cluster experiences a downtime until it is restored.
- With Hadoop 2.0 and Standby NameNodes, a mechanism for true high availability was realized.

# Hadoop Cluster Architecture with HA



- Standby NameNodes provide automatic failover if primary NameNode fails.
- Achieving high availability with Standby NameNodes requires shared storage between the primary and standbys (for the edit log).
- Though there are two options for the necessary shared storage - **NFS** and **Quorum Journal Manager(QJM)** - only QJM is considered production-ready.

## Name Node and QJM

- The QJM is a dedicated HDFS implementation, designed for the sole purpose of providing a highly available edit log, and is the recommended choice for most HDFS installations.
- Using QJM to maintain consistency of Active and Standby state requires that both nodes be able to communicate with a group of JournalNodes (JNs).
- When the Active node modifies the namespace, it logs a record of the change to a majority of JournalNodes. The StandbyNode watches the JNs for changes to the edit log and applies them to its own namespace.

# Working with HDFS

# Basic HDFS Commands



## Help

```
$ hadoop fs -help  
$ hdfs dfs -help
```

## Basic File Operations

\$ <b>hadoop fs</b> -mkdir <i>mydir</i>	Create directory
\$ <b>hadoop fs</b> -rmdir <i>mydir</i>	Remove a directory
\$ <b>hadoop fs</b> -ls	List files
\$ <b>hadoop fs</b> -ls <i>specific_dir</i>	List files in a dir
\$ <b>hadoop fs</b> -ls -R	Recursive listing
\$ <b>hadoop fs</b> -cat <i>myfile</i>	Read the file content
\$ <b>hadoop fs</b> -cp <source> <dist>	Copy file to a dir with in Hadoop
\$ <b>hadoop fs</b> -mv <source> <dist>	Move a file within hadoop

# Basic HDFS Commands



## Basic File Operations

<code>\$ hadoop fs -rm myfile</code>	Delete a file from Hadoop
<code>\$ hadoop fs -rm -R mydir</code>	Delete files recursively
<code>\$ hadoop fs -chmod 777 myfile</code>	Change permission on a file/dir
<code>\$ hadoop fs -chmod -R 777 mydir</code>	Change permissions recursively

## Moving files between Linux & HDFS

<code>\$ hadoop fs -copyFromLocal &lt;src&gt; &lt;dist&gt;</code>	Copy from Linux to HDFS
<code>\$ hadoop fs -copyToLocal &lt;src&gt; &lt;dist&gt;</code>	Copy from HDFS to Linux
<code>\$ hadoop fs -get &lt;src&gt; &lt;dist&gt;</code>	Copy from Linux to HDFS
<code>\$ hadoop fs -put &lt;src&gt; &lt;dist&gt;</code>	Copy from HDFS to Linux
<code>\$ hadoop fs -moveFromLocal &lt;src&gt; &lt;dist&gt;</code>	Move from Linux to HDFS
<code>\$ <b>hadoop distcp -update hdfs://192.168.1.1/xxx hdfs://192.168.1.2/yyy</b></code>	

