



Apache HBASE

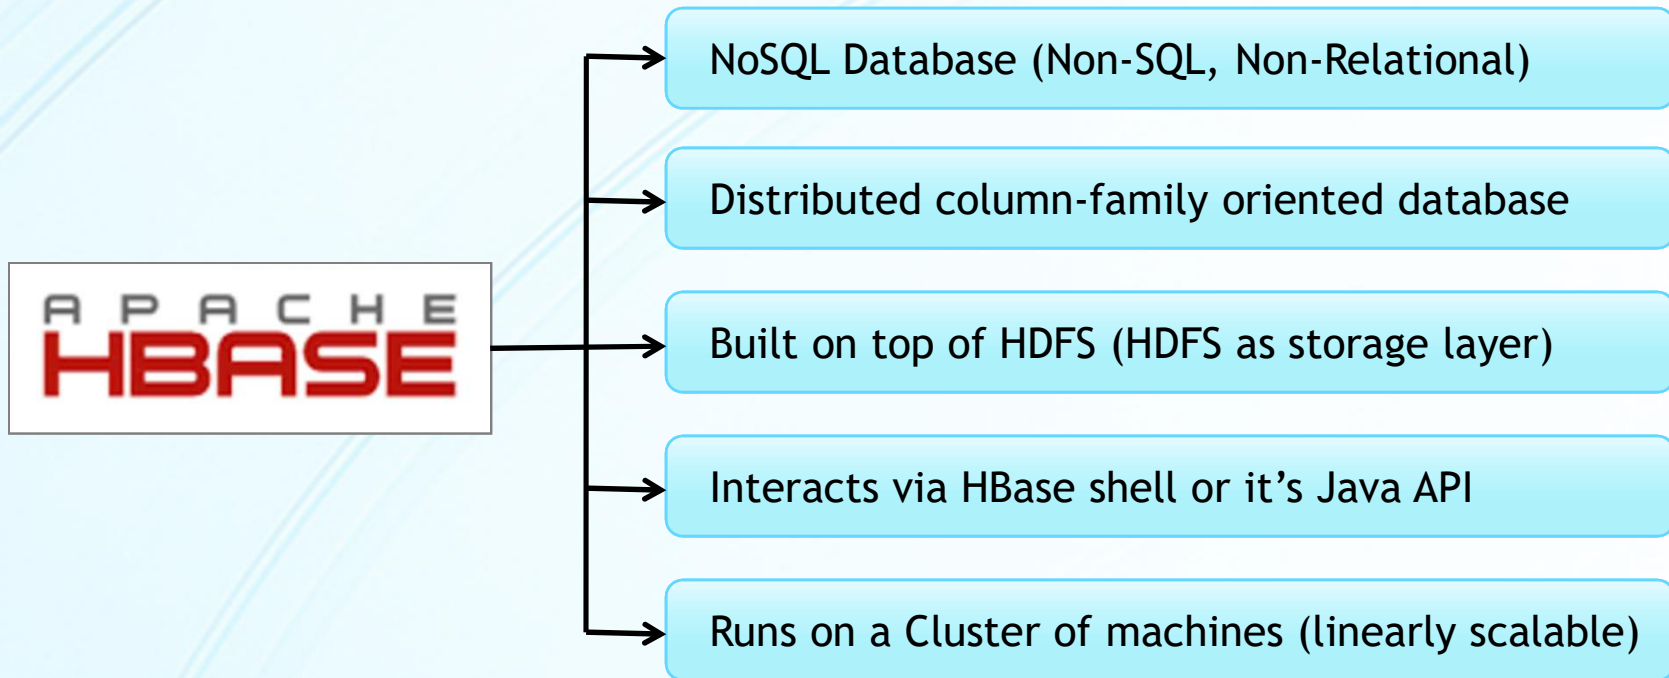


Agenda

- ✓ What is HBase?
- ✓ HBase Features & Use-cases
- ✓ HBase Run Modes
- ✓ Getting started with HBase
- ✓ HBase Data Model
- ✓ HBase Commands
- ✓ HBase Architecture



What is HBase?



HBase Features

APACHE
HBASE



Highly Scalable

Slaves nodes can be added on the fly to increase scalability.



HBase Features

The logo for Apache HBase, with "APACHE" in a grey, spaced-out font above "HBASE" in a bold, red font.A diagram showing the Apache HBase logo in a box. A vertical line descends from the box, with two horizontal arrows pointing to two light blue rounded rectangular boxes. The first box contains the text "Highly Scalable" and the second box contains the text "Dynamic Schema".

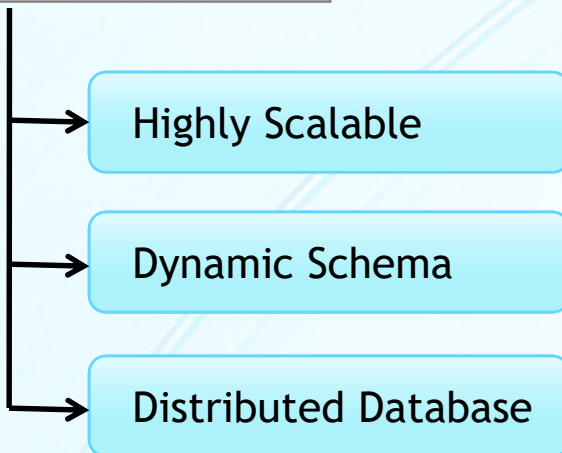
Highly Scalable

Dynamic Schema

Columns can be added on the fly. Columns are not part of the table schema.



HBase Features

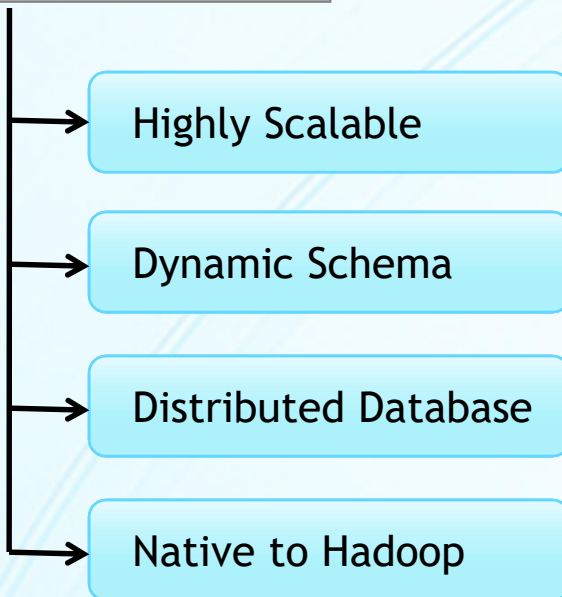


Data gets partitioned automatically as the data grows. HBase provides “automatic sharding” - tables are dynamically distributed by the system to different region servers when they become too large.



HBase Features

APACHE
HBASE



Can utilize MR framework and can use HDFS for storage there by gets all the advantages of Hadoop.



HBase Use Cases


Real time analytics



Search Engines



Storing RSS feeds



Messaging Infrastructure



HBase Run Modes



Stand-alone Mode

Default mode

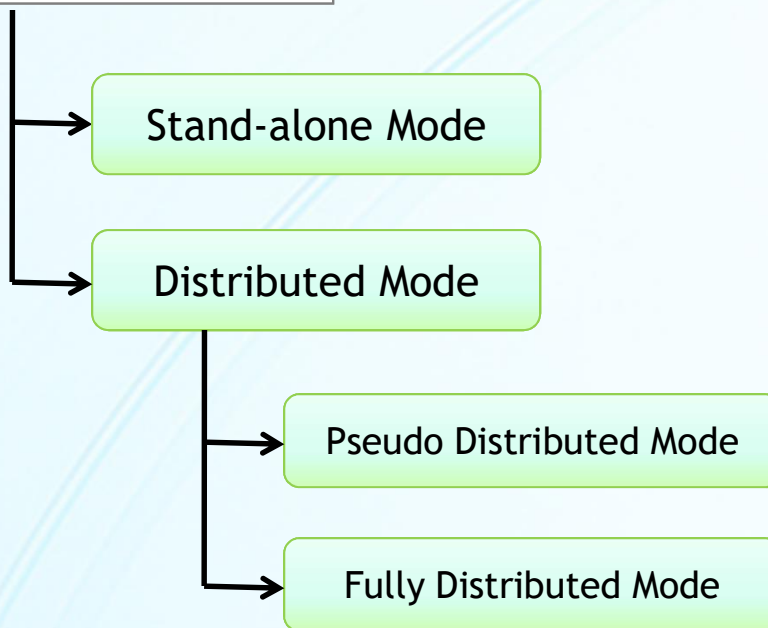
Does not use HDFS — uses local file system

Runs all HBase daemons and a local ZooKeeper in the same JVM process

ZooKeeper binds to a well-known port so that clients may talk to HBase.



HBase Run Modes




All daemons run on a single node

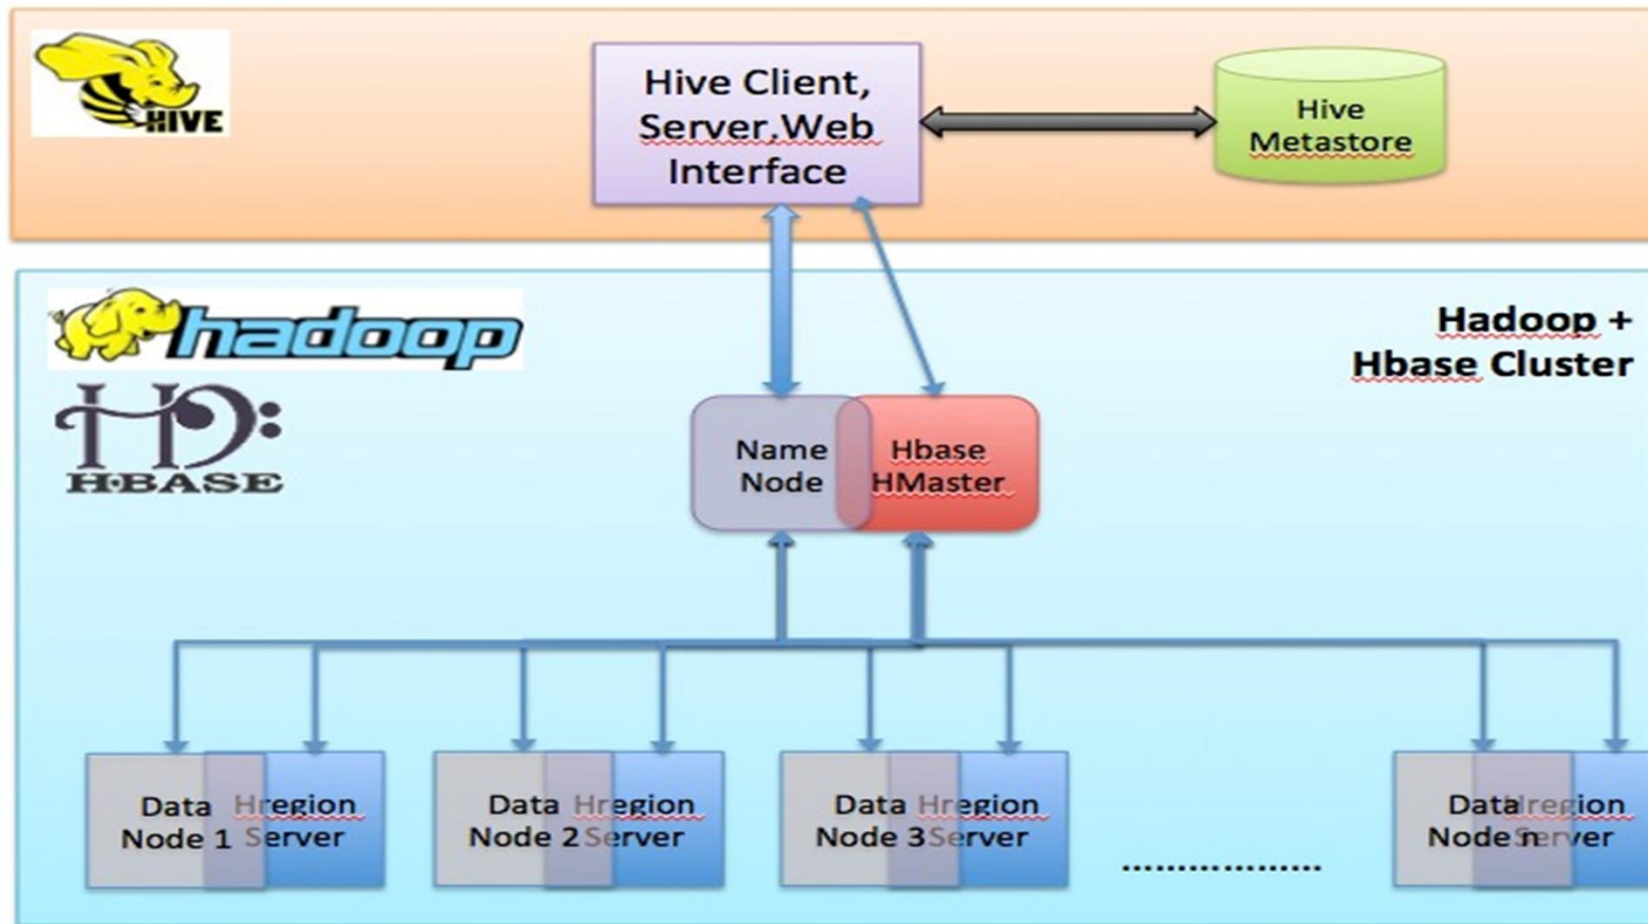
The daemons are spread across multiple, physical servers in the cluster.



HBase Architecture



HBase in Hadoop Ecosystem



Installing & Running HBase

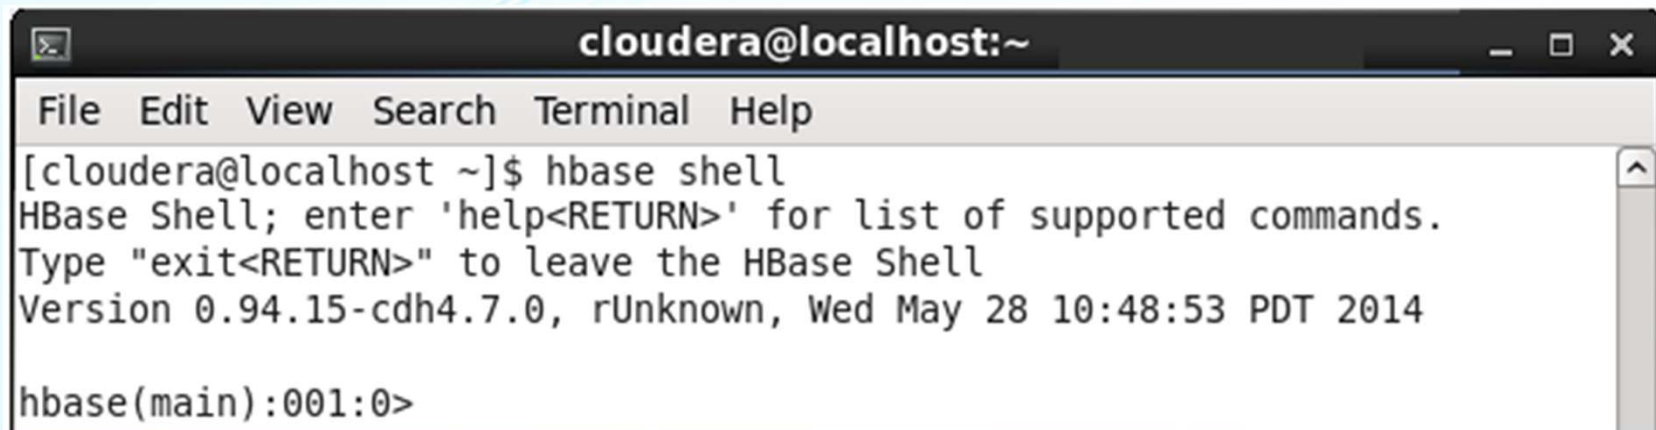
- HBase installs and runs on top of Hadoop. Hadoop needs to be installed and running before you install HBase.
- Once HBase is installed using either tar-balls or RPM packages, the main config file that has to be edited is `hbase-site.xml`
- The HBase `HMaster` process typically runs on the same machine where the Hadoop NameNode is installed
- The HBase `HRegionServer` processes run on each of the DataNodes along with the Hadoop DataNode and NodeManager processes



Invoking HBase Shell

You can enter into HBase Shell using the following command:

```
$hbase shell
```

A terminal window titled 'cloudera@localhost:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal shows the command '[cloudera@localhost ~]\$ hbase shell' being executed. The output is 'HBase Shell; enter \'help<RETURN>\' for list of supported commands. Type "exit<RETURN>" to leave the HBase Shell' followed by 'Version 0.94.15-cdh4.7.0, rUnknown, Wed May 28 10:48:53 PDT 2014'. The prompt 'hbase(main):001:0>' is shown at the bottom.

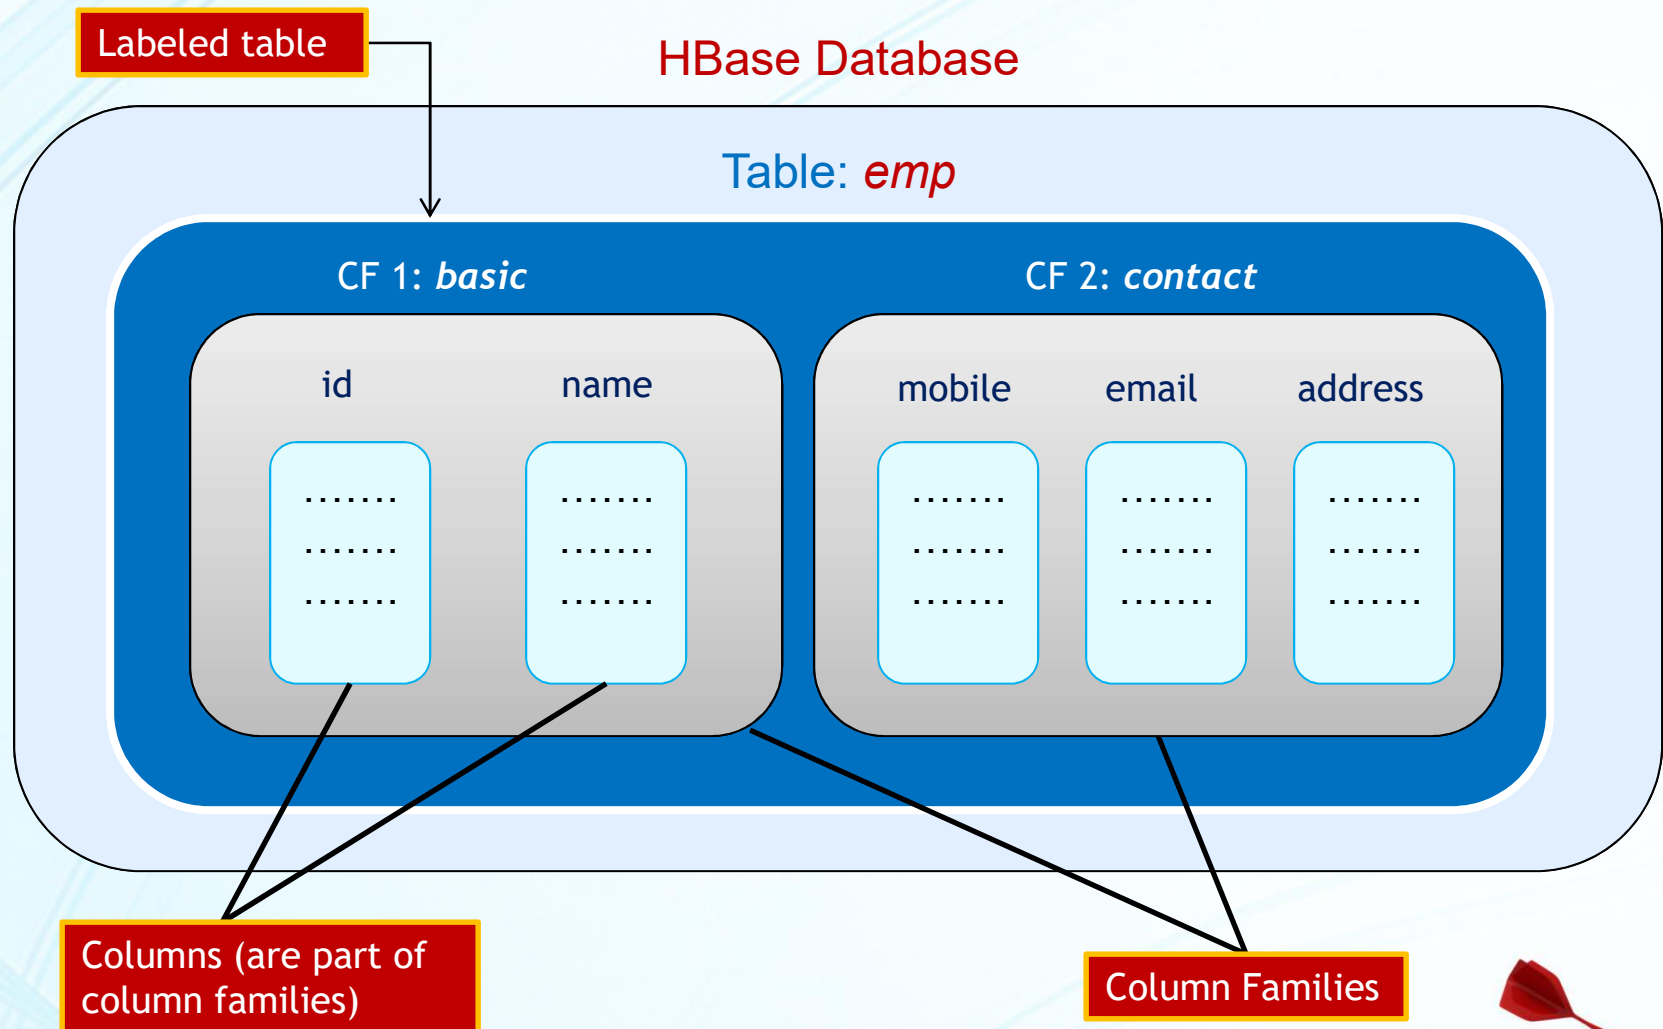
```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hbase shell  
HBase Shell; enter 'help<RETURN>' for list of supported commands.  
Type "exit<RETURN>" to leave the HBase Shell  
Version 0.94.15-cdh4.7.0, rUnknown, Wed May 28 10:48:53 PDT 2014  
hbase(main):001:0>
```



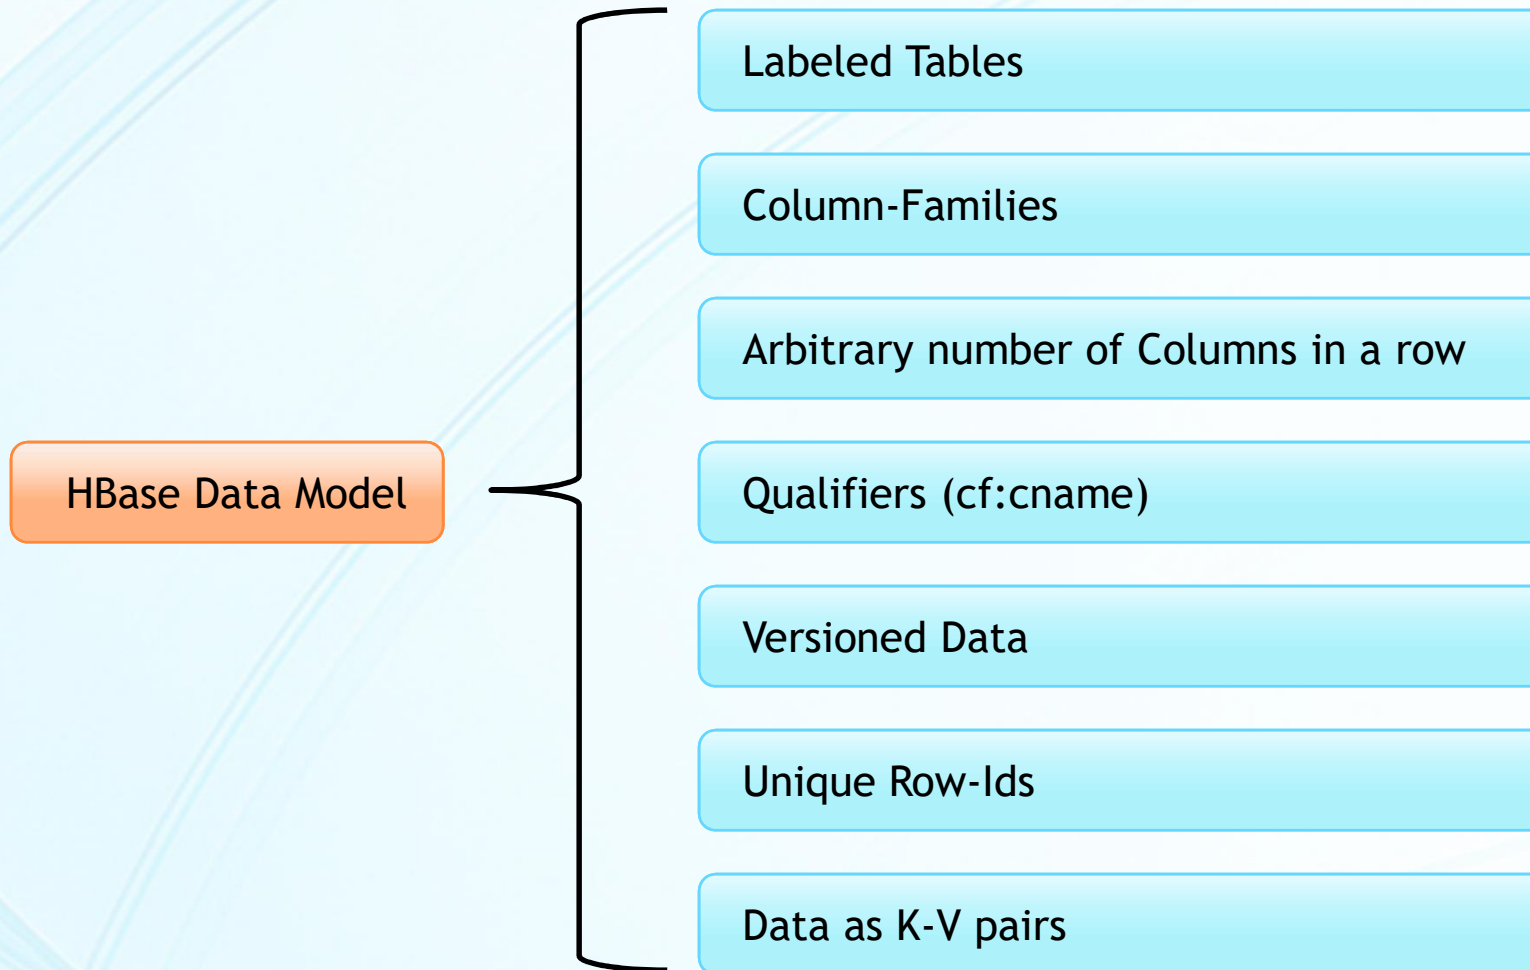
Let's understand how data is organized in HBase



HBase Data Model



HBase Data Model



HBase Data Model

In HBase, data is stored in **labeled tables**. A table definition consists of the **table name** (label) and a set of **column-families**.

A column-family may have *any number of columns*.

Columns are NOT part of table definition. Columns are assigned at the time of data insertion and *every row may not have the same number of columns*.

A table's column families must be specified up front as part of the table schema, but new column family members can be added on demand (using 'alter' command).

All column family members *have a common prefix*. For example, the columns *basic:id* and *basic:name* are both members of the '*basic*' column family. In this example, '*id*' and '*name*' are called *qualifiers*.



HBase Data Model

Table cells are versioned. By default, their version is a timestamp auto-assigned by HBase at the time of cell insertion.

Table data is stored as **key-value pairs**. Both the key and the value are byte-arrays.

Every row should contain a **unique row-key**. Data gets automatically sorted based on row-key.

HBase automatically adds a timestamp to track the versions of the data.



Let's look at HBase Shell Commands



Create a HBase table

```
create 'emp', 'basic', 'contact', 'salary', 'skills'

// table-name: t1, family: f1, versions to maintain: 4
create 't1', {NAME => 'f1', VERSIONS => 4}
create 't1', {NAME => 'f1', VERSIONS => 4, TTL => 2592000}

// table-name: t1, families: f1, f2, f3
create 't1', {NAME => 'f1'}, {NAME => 'f2'}, {NAME => 'f3'}

create 't1', 'f1', 'f2', 'f3' // same as above
```



Alter a table

// alter or add a column-family

```
alter 't1', NAME => 'f1', VERSIONS => 5
```

// delete a column-family

```
alter 't1', NAME => 'f3', METHOD => 'delete'
```



Drop a table

```
disable 'emp'  
drop 'emp'
```

NOTE: to drop, you need to disable the table



Insert data in a Table

```
put 'emp', 'r001', 'basic:empid', '100'  
put 'emp', 'r001', 'basic:empname', 'Raju'  
put 'emp', 'r001', 'basic:deptid', '1'  
put 'emp', 'r001', 'contact:email', 'raju@gmail.com'  
put 'emp', 'r001', 'contact:phone', '9876543210'
```

Table Name

Row Id

Qualified Key

Value



Scan a table

Scan command scans the entire table and gets you the data of the entire table.

```
scan 'emp`  
scan 'emp', {COLUMNS=>['basic:name', 'basic:id']}  
scan 'emp', {COLUMNS => ['id', 'name'], LIMIT => 10, STARTROW => 'r0003'}  
scan 'emp', {COLUMNS => 'id', TIMERANGE => [1303668804, 1303668904]}
```

```
hbase(main):038:0> scan 'empinfo'  
ROW                                COLUMN+CELL  
100ABCD                            column=employee:firstName, timestamp=1545966262094, value=Veer4  
100_Veer                            column=dept:dname, timestamp=1545966108929, value=Operations  
100_Veer                            column=employee:DOB, timestamp=1545966081267, value=12-12-2012  
100_Veer                            column=employee:firstName, timestamp=1545965933625, value=Veer  
100_Veer                            column=employee:lastName, timestamp=1545966053288, value=Nagaraju  
100_seshu                          column=employee:hd, timestamp=1545966316687, value=12-12-2010  
100_seshu                          column=employee:lastName, timestamp=1545966331643, value=xxxxxxx  
3 row(s) in 0.0460 seconds
```



Query data from a table

```
get 't1', 'r1'  
get 't1', 'r1', {TIMERANGE => [ts1, ts2]}  
get 't1', 'r1', {COLUMN => 'cf1:c1'}  
get 't1', 'r1', {COLUMN => ['cf1:c1', 'cf1:c2', 'cf1:c3']}  
get 't1', 'r1', {COLUMN => 'cf1:c1', TIMESTAMP => ts1}  
get 't1', 'r1', {COLUMN => 'cf1:c1', VERSIONS => 4}  
get 't1', 'r1', 'cf1:c1'  
get 't1', 'r1', ['cf1:c1', 'cf1:c2']
```

Table Name

Row Id

Column list

Column Options



Delete a cell

```
delete 'emp', 'r001', 'basic:name'
```

Table Name

Row Id

Column list



Describe the table

```
describe 'emp'
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
hbase(main):006:0> describe 'empinfo'  
DESCRIPTION  
'empinfo', {NAME => 'address', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0', VERSIONS => '3', COMPRESSION => 'NONE', MIN_VERSIONS => '0', TTL => '2147483647', KEEP_DELETED_CELLS => 'false', BLOCKSIZE => '65536', IN_MEMORY => 'false', ENCODE_ON_DISK => 'true', BLOCKCACHE => 'true'}, {NAME => 'dept', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0', VERSIONS => '3', COMPRESSION => 'NONE', MIN_VERSIONS => '0', TTL => '2147483647', KEEP_DELETED_CELLS => 'false', BLOCKSIZE => '65536', IN_MEMORY => 'false', ENCODE_ON_DISK => 'true', BLOCKCACHE => 'true'}, {NAME => 'employee', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0', VERSIONS => '3', COMPRESSION => 'NONE', MIN_VERSIONS => '0', TTL => '2147483647', KEEP_DELETED_CELLS => 'false', BLOCKSIZE => '65536', IN_MEMORY => 'false', ENCODE_ON_DISK => 'true', BLOCKCACHE => 'true'}, {NAME => 'salgrade', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'NONE', REPLICATION_SCOPE => '0', VERSIONS => '3', COMPRESSION => 'NONE', MIN_VERSIONS => '0', TTL => '2147483647', KEEP_DELETED_CELLS => 'false', BLOCKSIZE => '65536', IN_MEMORY => 'false', ENCODE_ON_DISK => 'true', BLOCKCACHE => 'true'}  
1 row(s) in 0.0800 seconds
```



Executing a HBase Script

```
$hbase shell ./scriptfile.txt
```



The script file to be executed

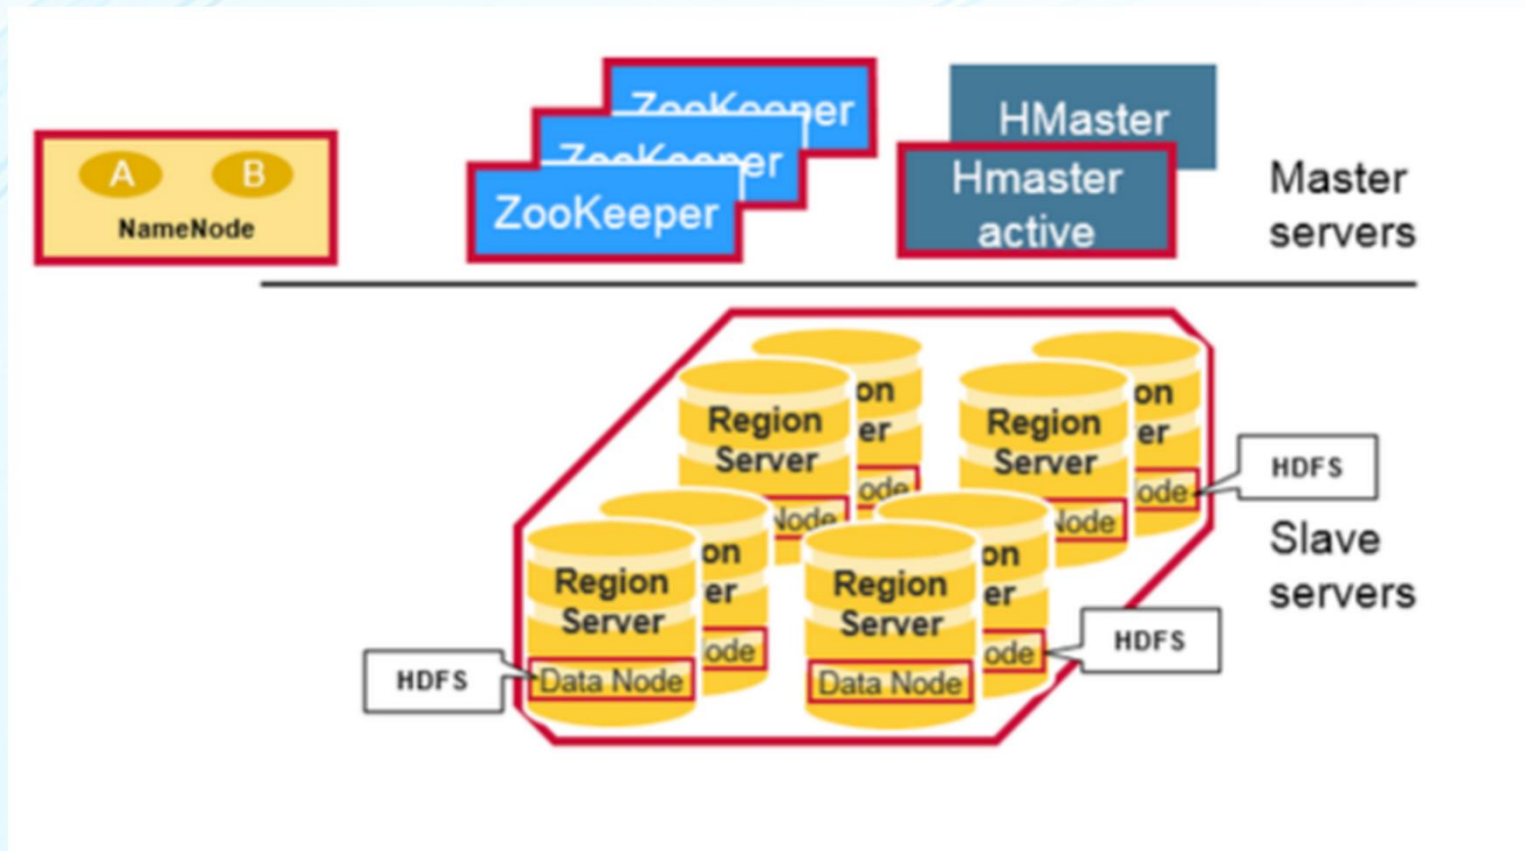
- It is more common to execute script files than running ad-hoc commands from the shell.
- You can write all the HBase commands in a single file and execute all the commands in that script file using the above command.



HBase Architecture



HBase Architecture



HBase Daemons: **HMaster** (master daemon) & **RegionServer** (slave daemon)

HBase Components

HBase follows master-slave architecture. HBase is made up of an HBase 'Master' node orchestrating a cluster of one or more 'RegionServer' workers.



HBase Master

Responsible for bootstrapping a virgin install

Assigns regions to registered region-servers

Responsible for recovering RegionServer failures

Is lightly loaded



HBase Components

HBase follows master-slave architecture. HBase is made up of an HBase 'Master' node orchestrating a cluster of one or more 'RegionServer' workers.



Carry zero or more regions

Field client read/write requests

Manage region splits, informing the HBase master about the new regions



HBase Components

HBase follows master-slave architecture. HBase is made up of an HBase 'Master' node orchestrating a cluster of one or more 'RegionServer' workers.

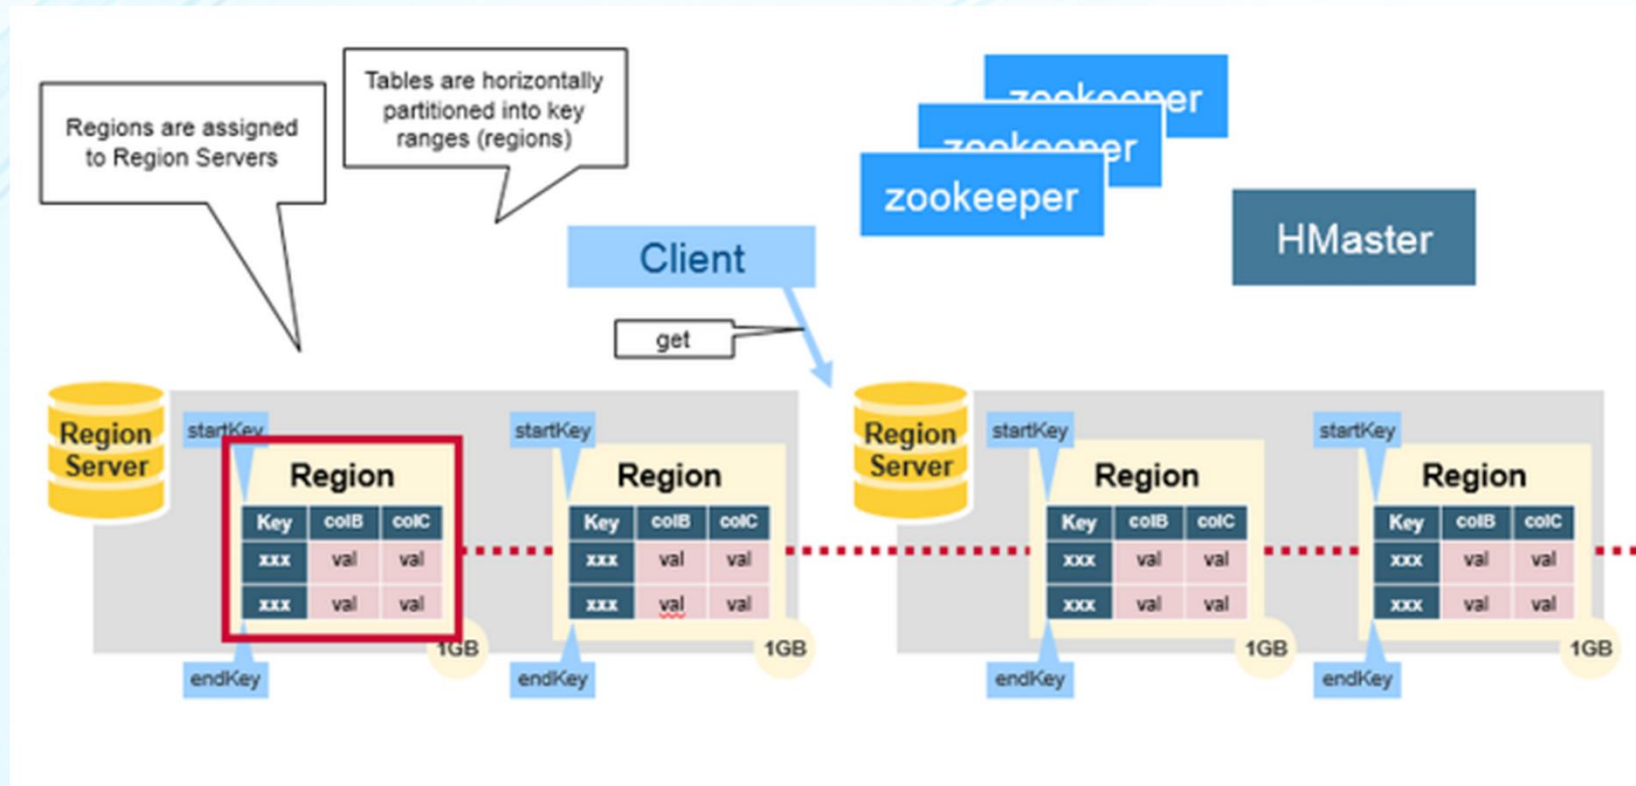


Authority on cluster state

Hosts vitals like location of the hbase:meta catalog table and the address of the current cluster master.

Assignment of regions is mediated via ZooKeeper

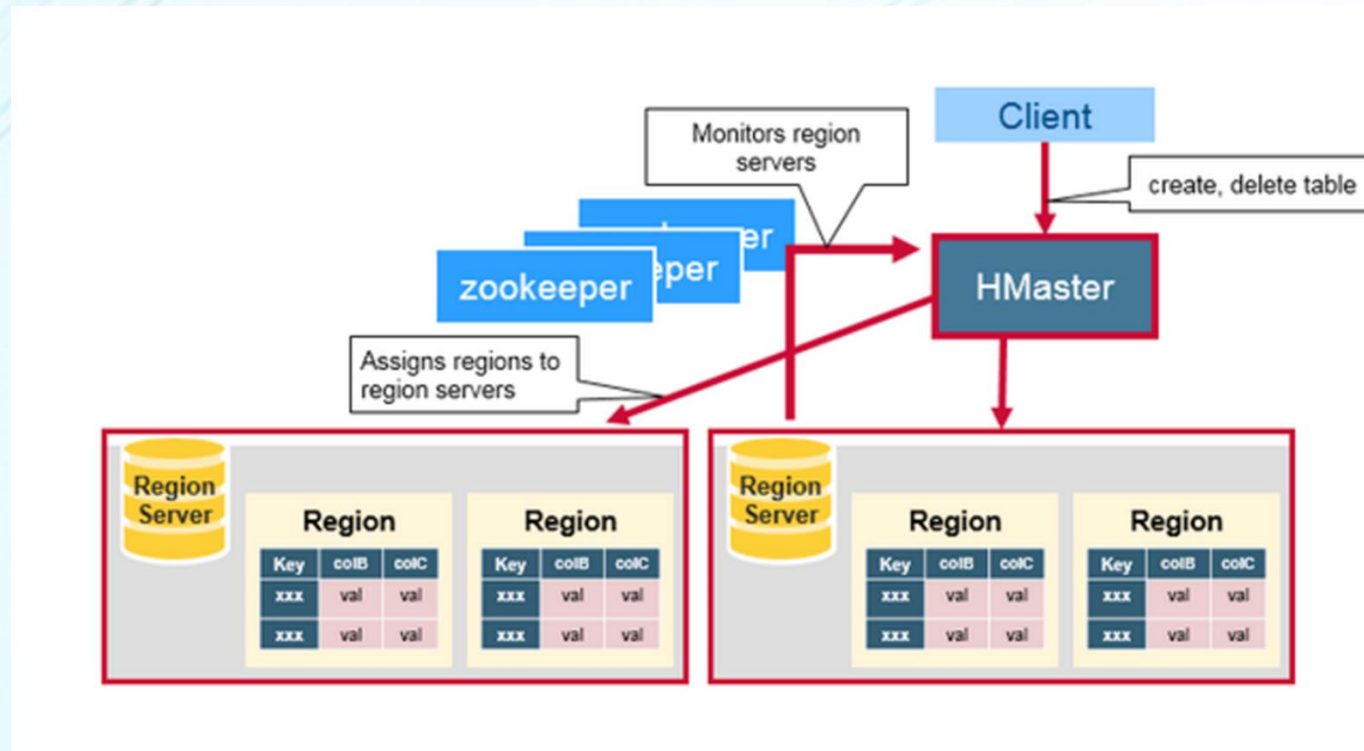
HBase Architecture contd..



- Tables are horizontally partitioned into key regions (regions)
- Regions are assigned to RegionServers



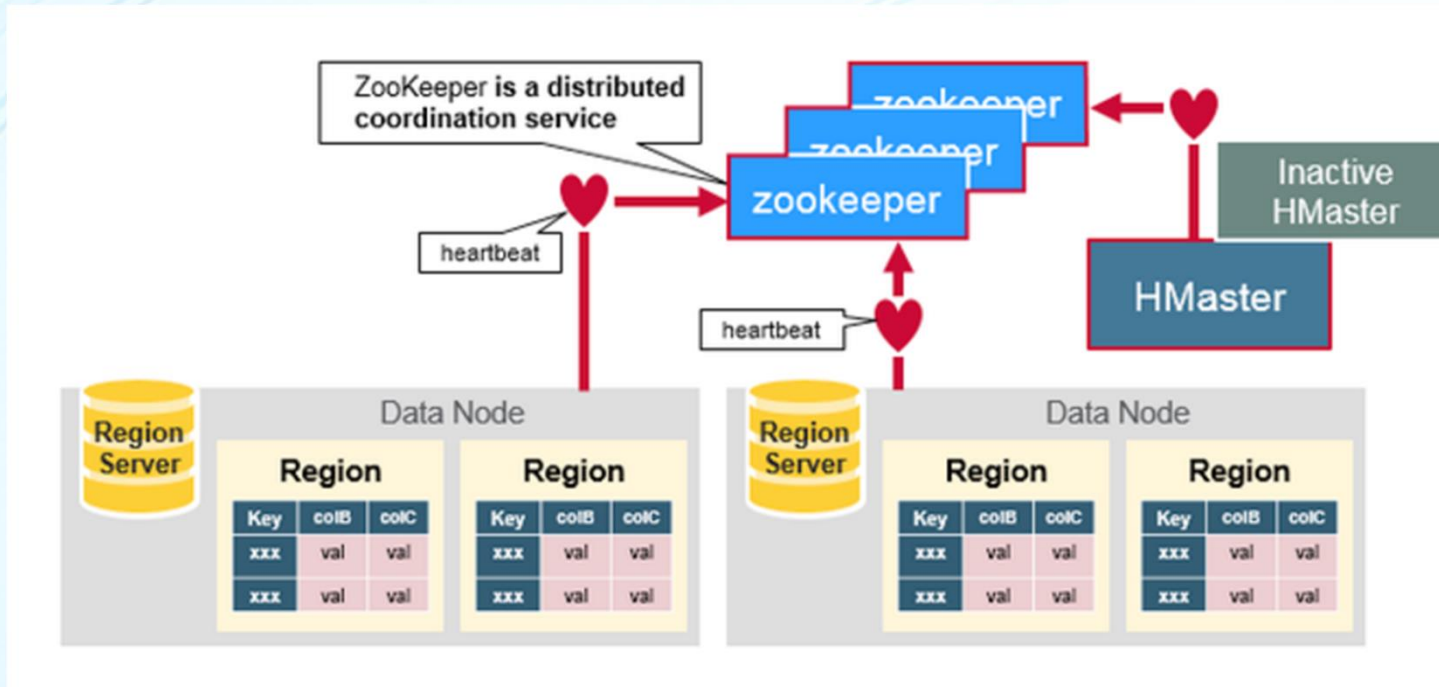
HBase Architecture contd..



- HMaster assigns regions to RegionServers
- HMaster monitors RegionServers



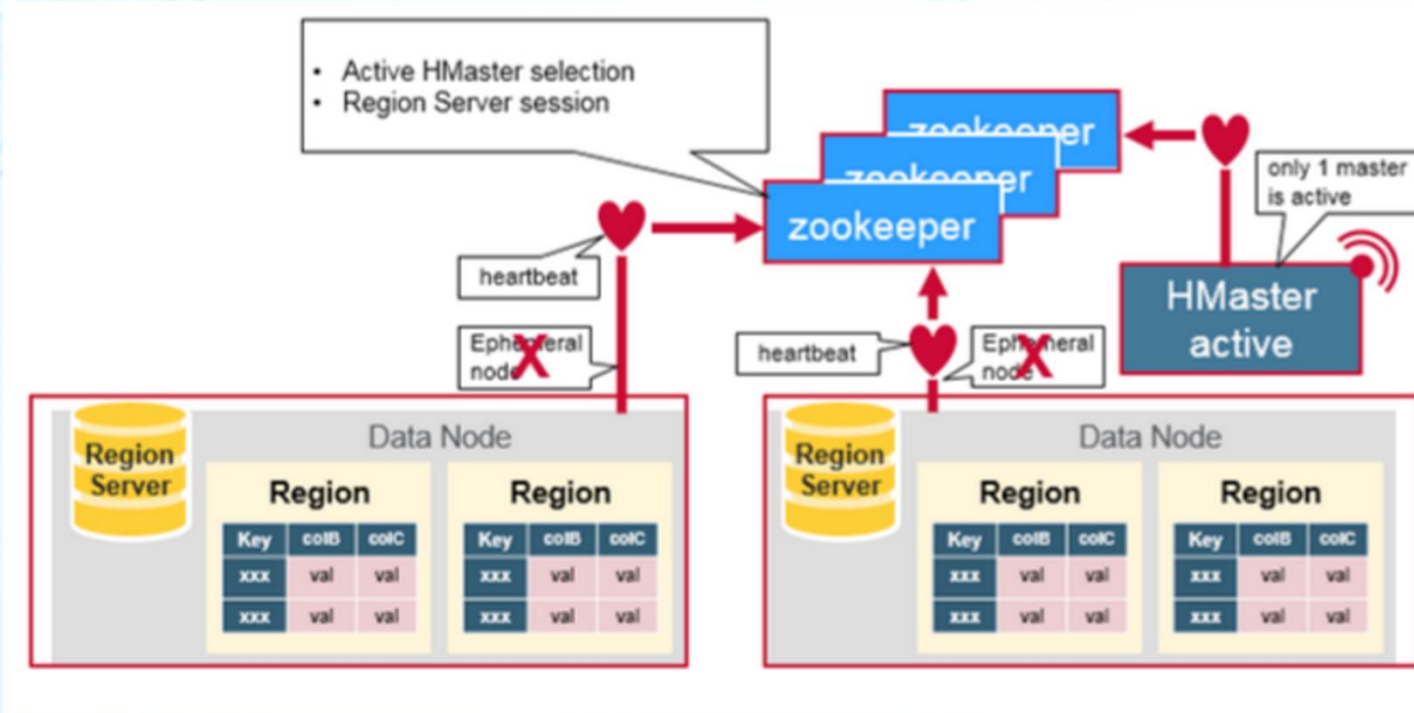
HBase Architecture contd..



- Zookeeper acts as distributed coordination service
- HMaster and RegionServers send heartbeats to ZooKeeper



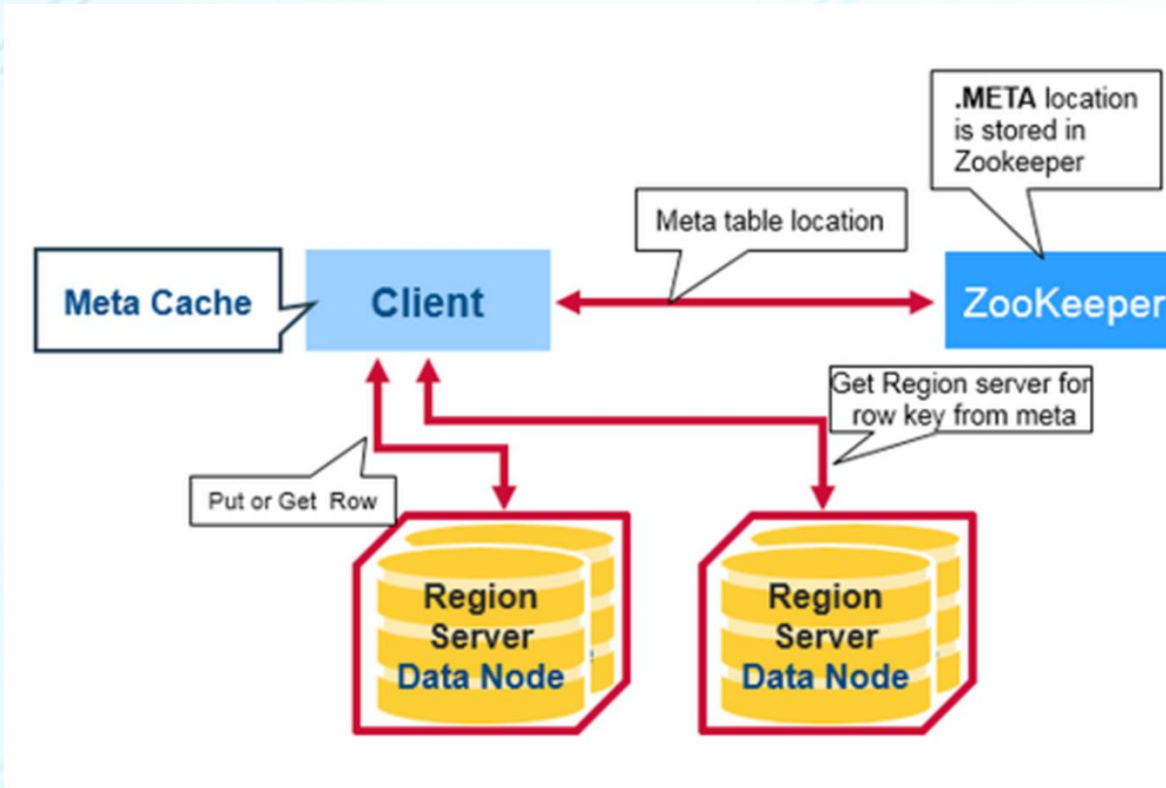
HBase Architecture contd..



- Zookeeper manages active HMaster selection and monitors RegionServer session.



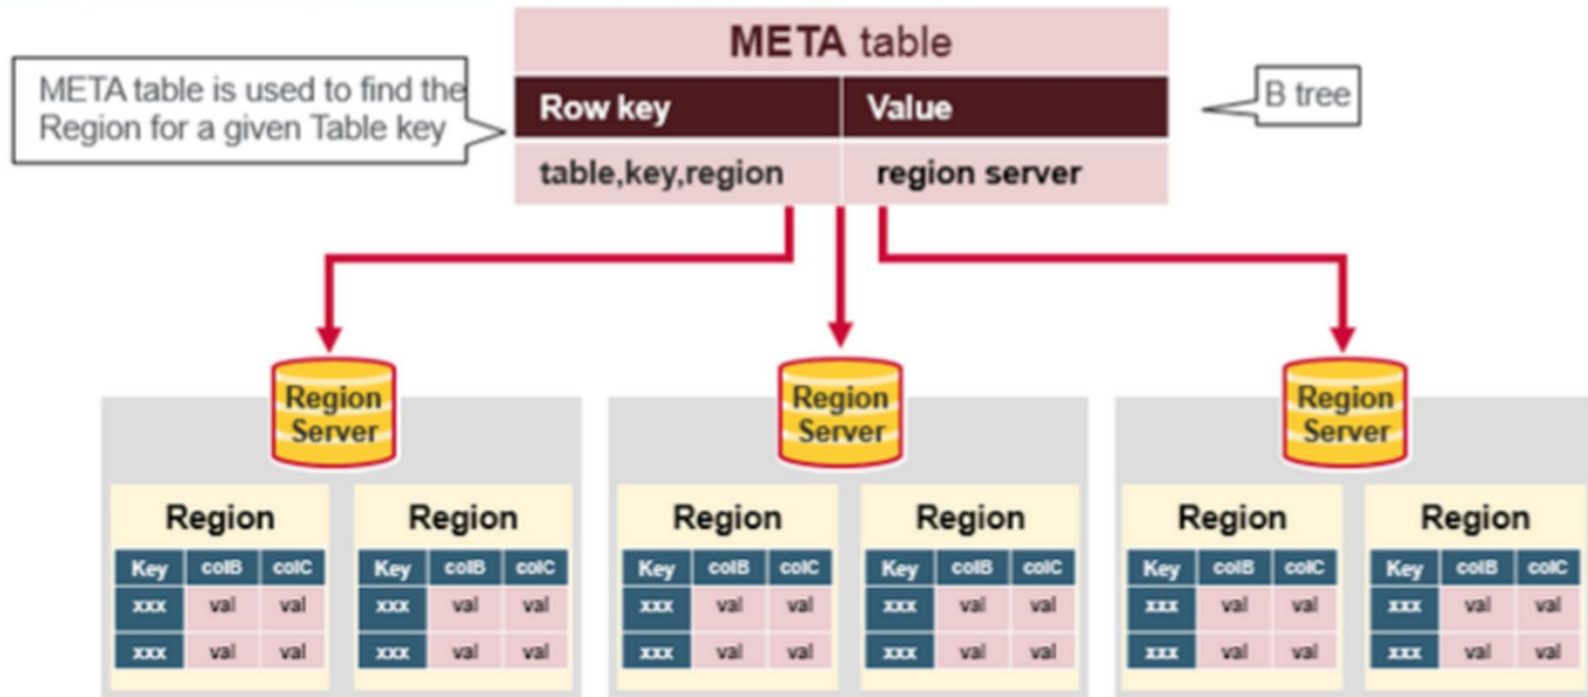
HBase Architecture contd..



Client receives meta info from ZK to decide which RegSrvr to read from or write to.



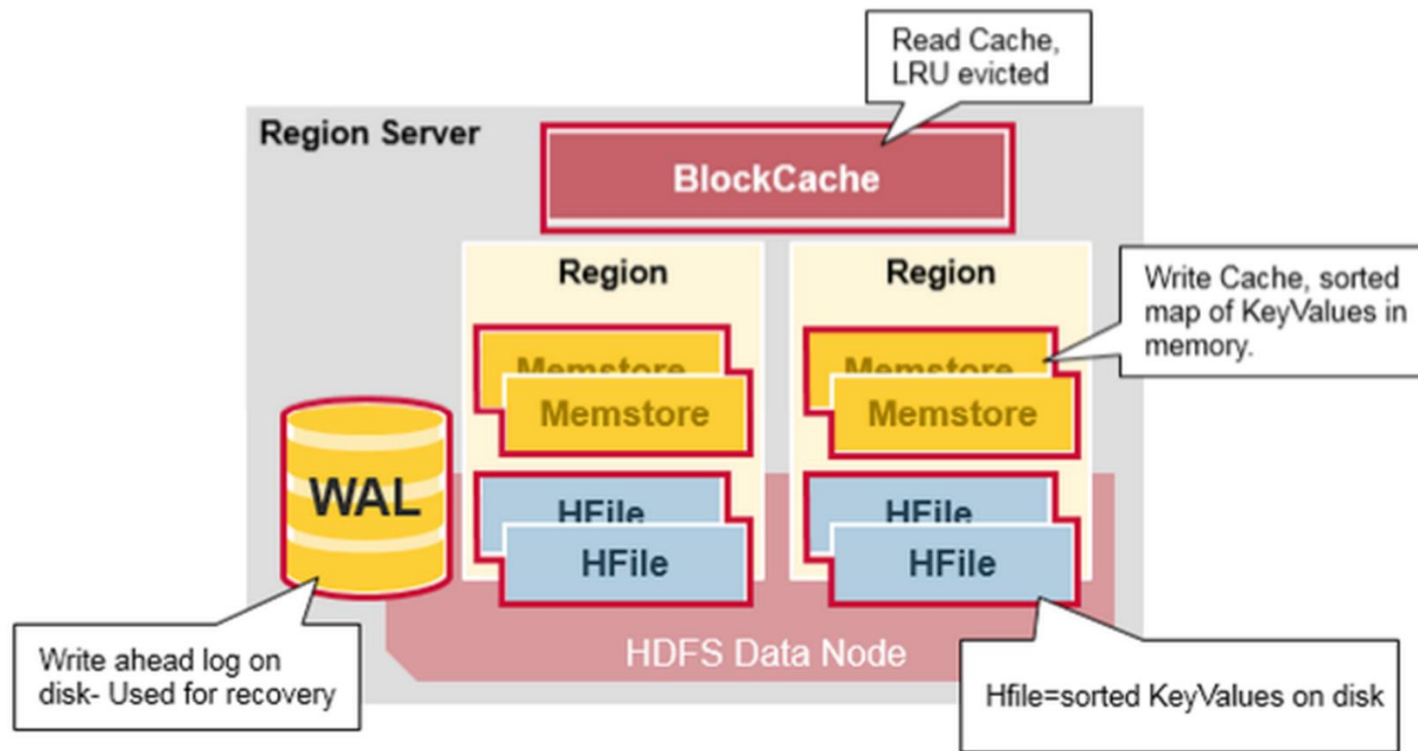
HBase Architecture contd..



Meta table is used to find the Region for a given table key.

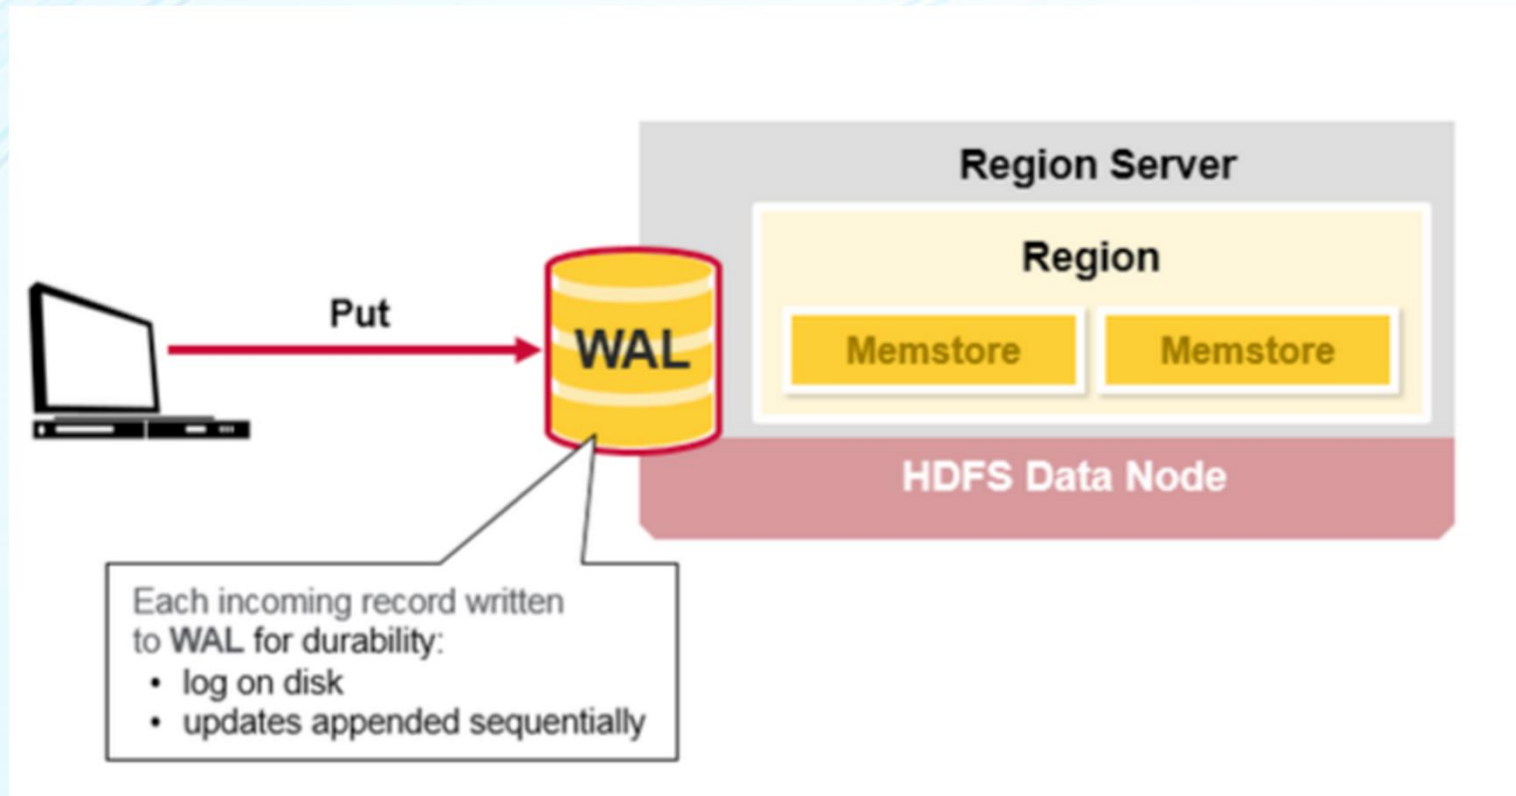


HBase Architecture contd..



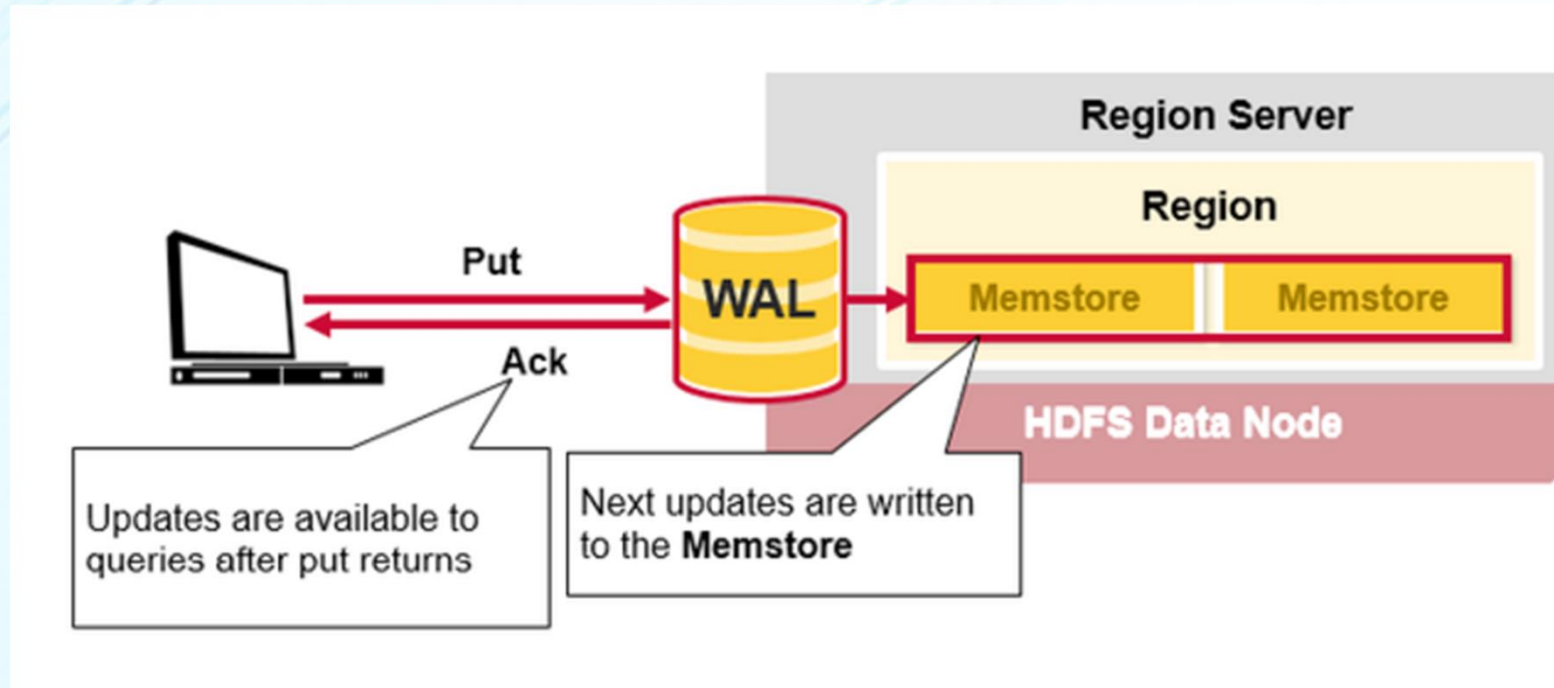
Writes arriving at a RegSrvr are first appended to a 'commit log' (WAL) and then added to an in-memory 'MemStore'. There is one MemStore for each column-family. That's why we say, column-family data gets sorted and stored at one place.

HBase Architecture contd..

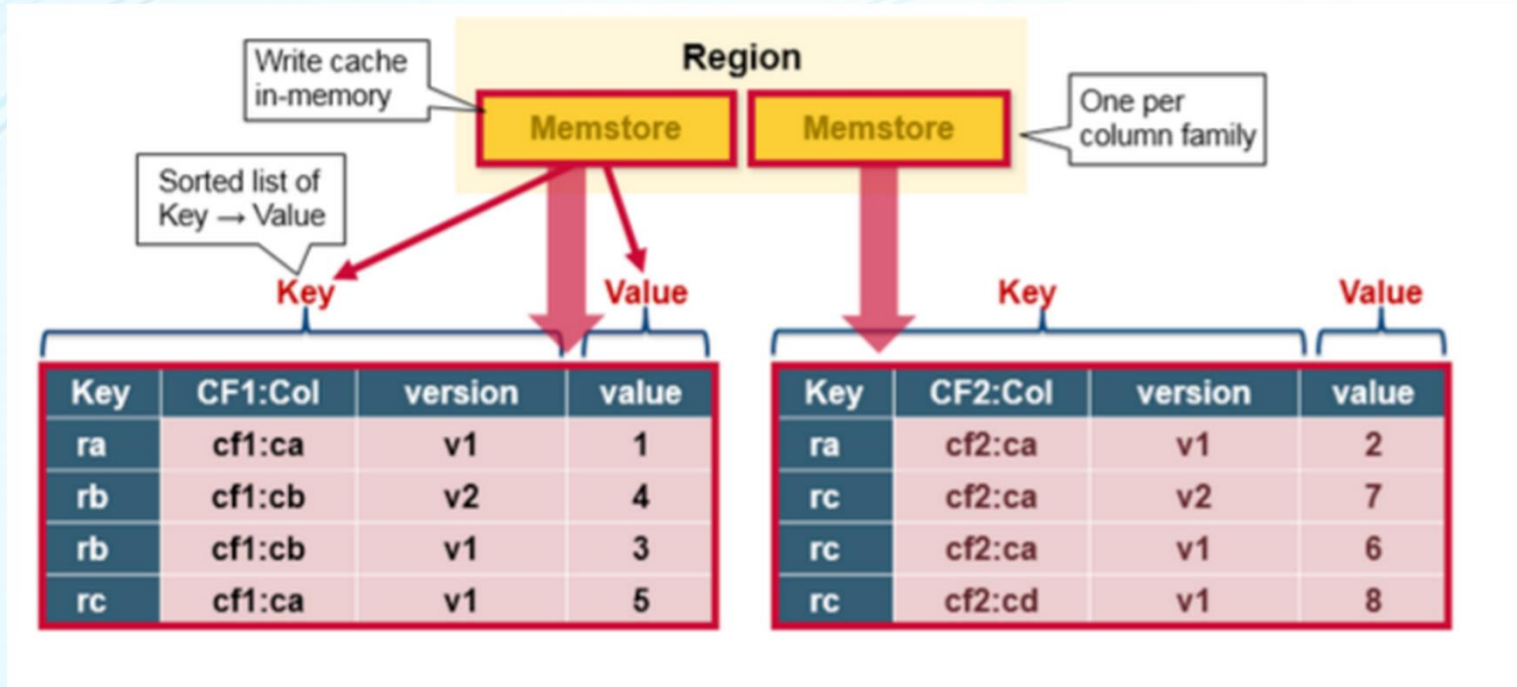


Each incoming record is written to WAL. WAL is stored on HDFS and updates are appended sequentially.

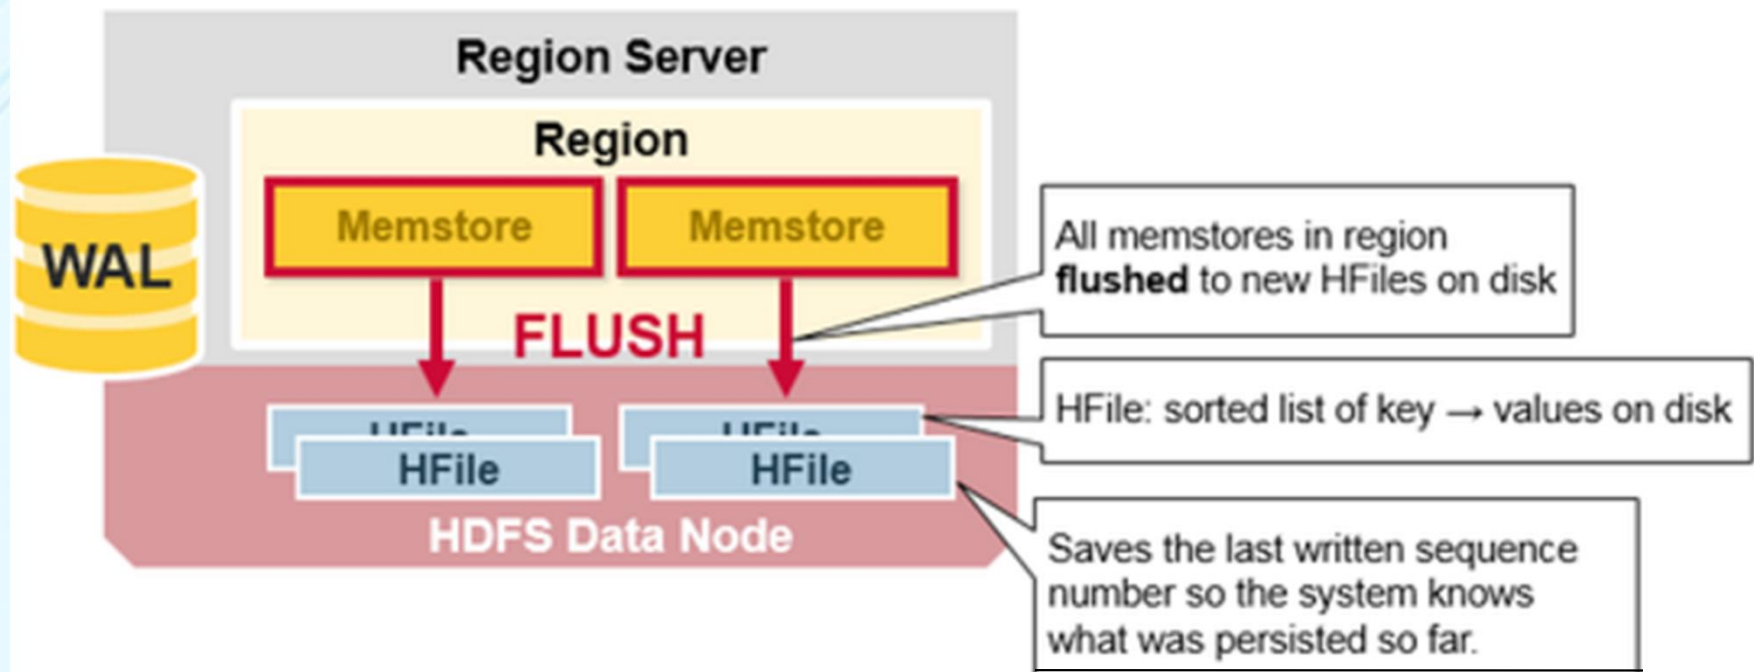
HBase Architecture contd..



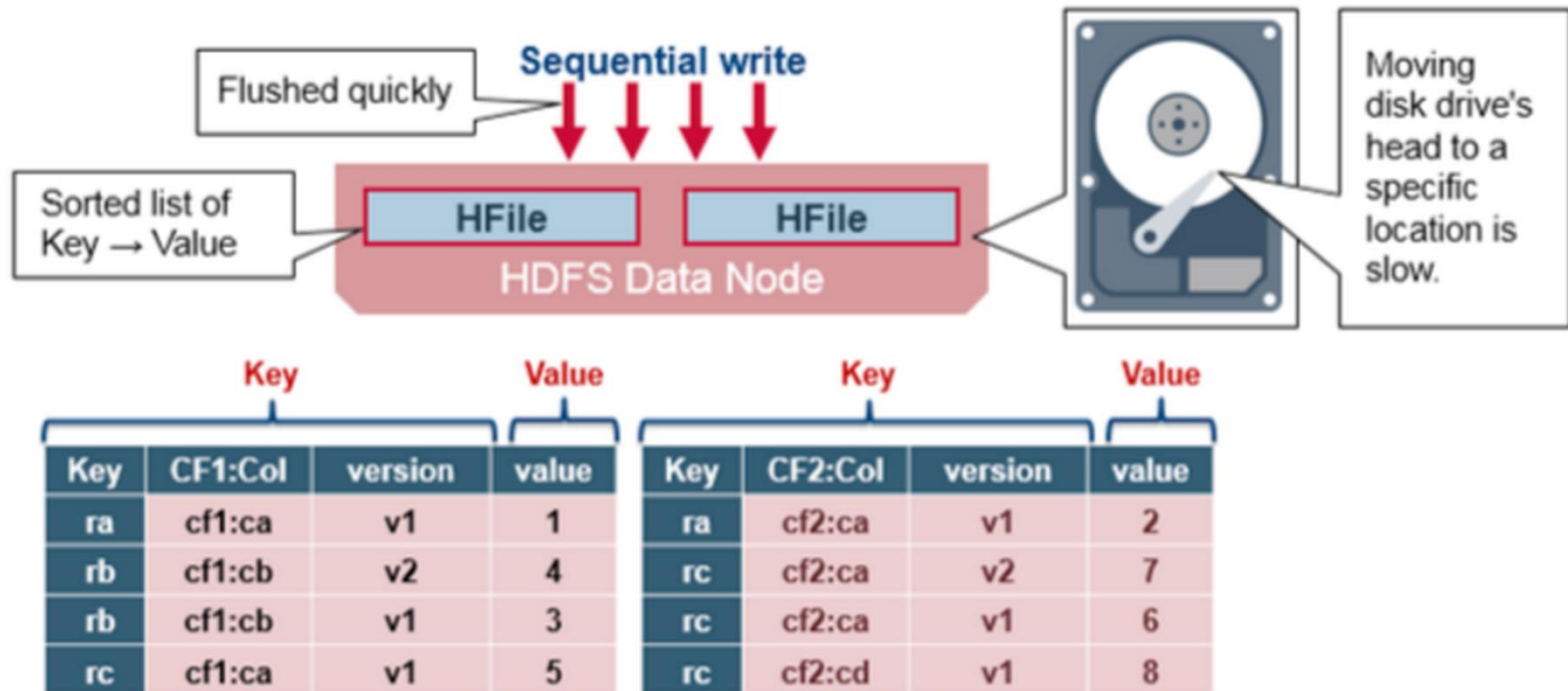
HBase Architecture contd..



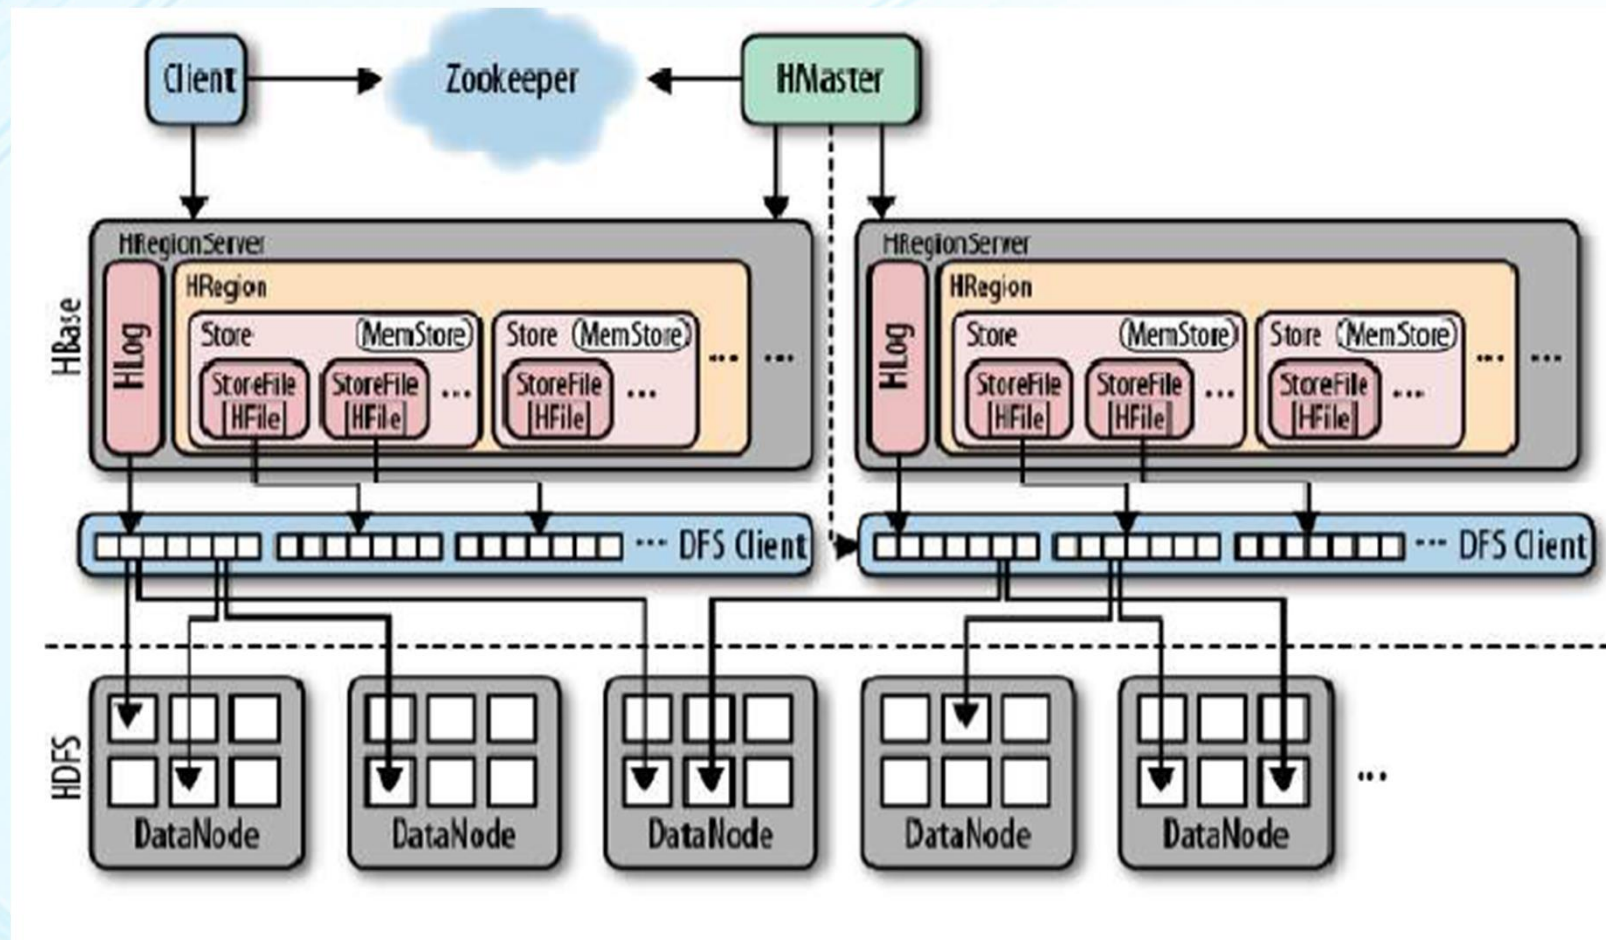
HBase Architecture contd..



HBase Architecture contd..



HBase Architecture contd..



Compaction

- The store files are monitored by a background thread to keep them under control.
- The flushes of MemStores (as HFiles) slowly build up an increasing number of on-disk files.
- If there are enough of them, the compaction process will combine them to a few, larger files.
- This goes on until the largest of these files exceeds the configured maximum store file size and triggers a region split.



THANK YOU

